



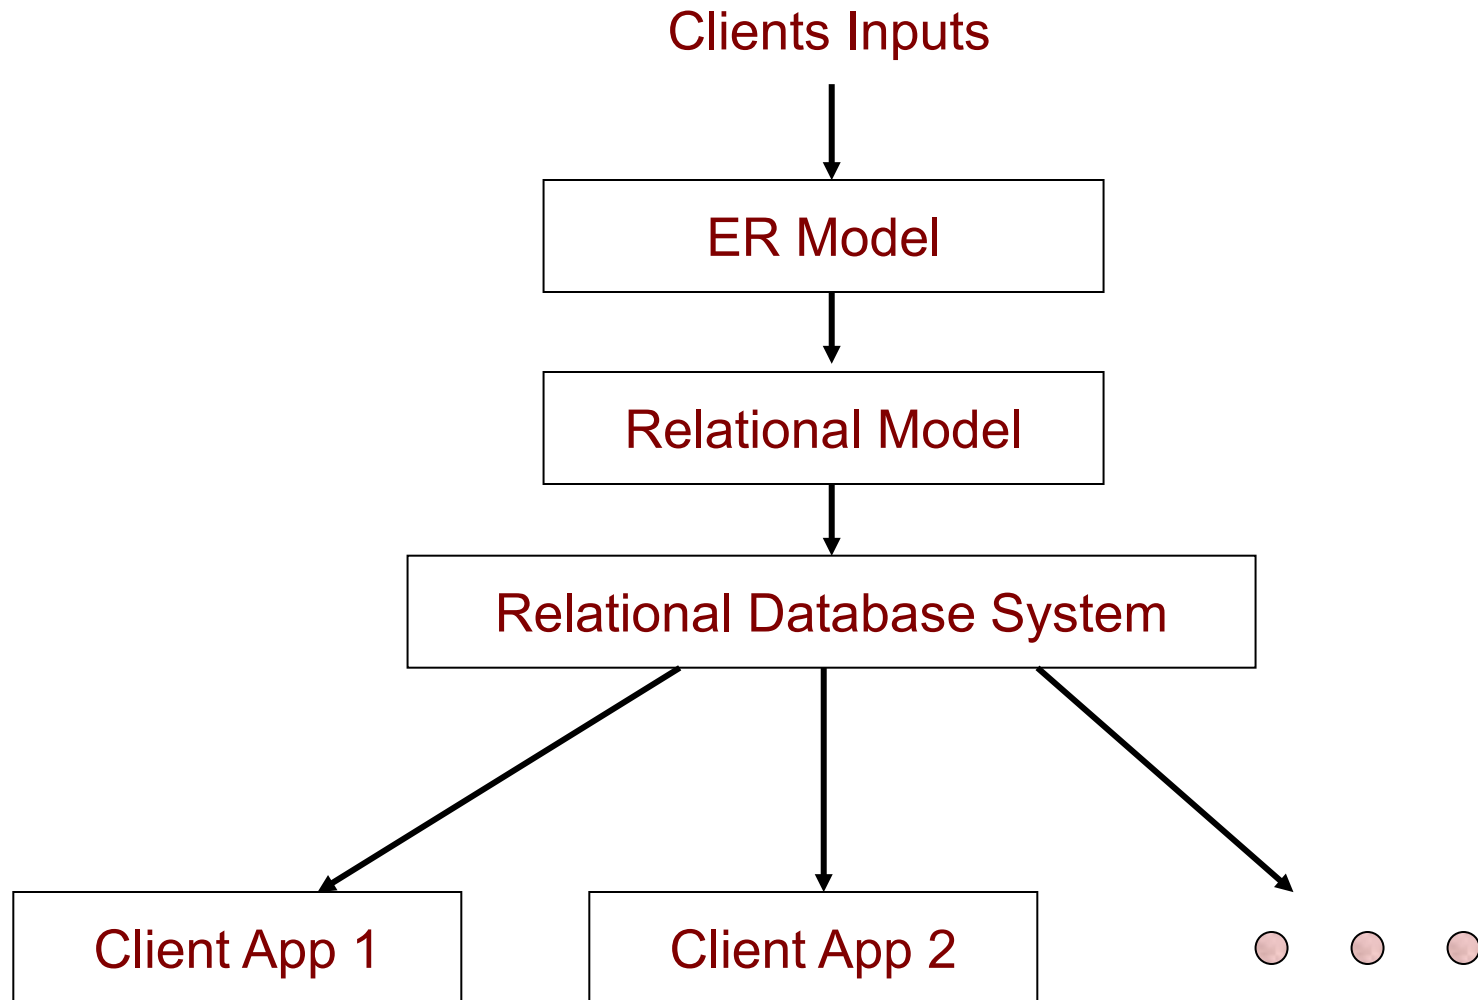
# **SE240: Introduction to Database Systems**

**Lecture 03:  
Relational Model**

# Outline

- **Relational Model**
  - **Relational Model Concepts**
  - Relational Database Schemas
- ER-to-Relational Mapping
  - Translating traditional ER diagrams
  - Translating Class Hierarchy

# Database Modeling



# Introduction

- The relational model was first introduced by Ted Codd of IBM Research in 1970
- Attracted due to its simplicity, elegance and mathematical foundations
- The model uses the concept of a mathematical relation – which looks like a **table** of values
- The **SQL Query Language** was developed by IBM in the 1970's

# Why Study the Relational Model?

- The relational model is by far the dominant data model
- It is the foundation for the leading DBMS products of many vendors such as Oracle, Microsoft SQL Server, etc.
- A major strength of the relational model is that it supports simple, powerful querying of data
- Queries can be written intuitively, and the DBMS can perform efficient evaluation

# The Relational Data Model

- The relational Model of Data is based on the concept of a *Relation*, a Relation is a mathematical concept based on the ideas of sets
- Represents data as a collection of relations
- **Table** of values
  - Row
    - Represents a collection of related data values
    - Fact that typically corresponds to a real-world entity or relationship
  - Table name and column names
    - Interpret the meaning of the values in each row *attribute*

Relation Name/Table Name

Attributes/Columns (collectively as a schema)

STUDENT			
Name	<u>Student-id</u>	Age	CGA
Chan Kin Ho	99223367	23	11.19
Lam Wai Kin	96882145	17	10.89
Man Ko Yee	96452165	22	8.75
Lee Chin Cheung	96154292	16	10.98
Alvin Lam	96520934	15	9.65

Tuples/Rows

$R(A_1, A_2, A_3, A_4)$

- **Relation**  $\leftrightarrow$  **table**; denoted by  $R(A_1, A_2, \dots, A_n)$  where  $R$  is a **relation name** and  $(A_1, A_2, \dots, A_n)$  is the **relation schema** of  $R$
- **Attribute** (column)  $\leftrightarrow$  denoted by  $A_i$   
Dom(Age): [0-100]  
Dom(EmpName): 50 alphabetic chars  
Dom(Salary): non-negative integer
- **Tuple (Record)**  $\leftrightarrow$  row
- **Attribute value**  $\leftrightarrow$  value stored in a table cell
- **Domain**  $\leftrightarrow$  legal type and range of values of an attribute, denoted by  $\text{dom}(A_i)$

# Relation Schema

- The Schema is a description of a Relation:
  - Denoted by  $R(A_1, A_2, \dots, A_n)$ ,  $R$  is the name of the relation
  - The attributes of the relation are  $A_1, A_2, \dots, A_n$
  - All legal values of an attribute is called its domain
- Example:
- STUDENT(Name, Student-id, Age, CGA)
  - STUDENT is the relation name
  - Defined over the four attributes: Name, Student-id, Age, CGA
  - A STUDENT state may include 5 STUDENTs; another 250 STUDENTs



# Relations States are Sets

- Relation State  $r$ , or  $r(R)$ : a specific state of relation  $R$  – this is a set of tuples (rows) *subset*
  - $r(R) = \{t_1, t_2, \dots, t_n\}$  where each  $t_i$  is an  $n$ -tuple,  $t_i = \langle v_1, v_2, \dots, v_n \rangle$  where each  $v_j$  element-of  $\text{dom}(A_j)$
  - **Mathematical relation** of degree  $n$  on the domains  $\text{dom}(A_1), \text{dom}(A_2), \dots, \text{dom}(A_n)$
  - **Subset** of the **Cartesian product** of the domains that define  $R$ :  $r(R) \subseteq \text{dom}(A_1) \times \text{dom}(A_2) \times \dots \times \text{dom}(A_n)$
- All tuples in a relation state  $r(R)$  form a set
- By definition, there cannot be duplicates in a set
- By definition, set elements (tuples) are **not ordered**, even though tuples frequently appear to be in the tabular form

# Example

- Let  $R(A_1, A_2)$  be a relation schema:
  - Let  $\text{dom}(A_1) = \{ 0, 1 \}$
  - Let  $\text{dom}(A_2) = \{ a, b, c \}$
- Then:  $\text{dom}(A_1) \times \text{dom}(A_2)$  is all possible combinations:  $\{ \langle 0, a \rangle, \langle 0, b \rangle, \langle 0, c \rangle, \langle 1, a \rangle, \langle 1, b \rangle, \langle 1, c \rangle \}$
- The relation state  $r(R) \subseteq \text{dom}(A_1) \times \text{dom}(A_2)$
- For example:  $r(R)$  could be  $\{ \langle 0, a \rangle, \langle 0, b \rangle, \langle 1, c \rangle \}$ 
  - this is one possible state  $r$  of the relation  $R$ , defined over  $A_1$  and  $A_2$
  - it has three 2-tuples:  $\langle 0, a \rangle, \langle 0, b \rangle, \langle 1, c \rangle$

# Values in Relations

- We refer to **component values** of a tuple  $t$  by:
  - $t[A_i]$  or  $t.A_i$
  - This is the value  $v_i$  of attribute  $A_i$  for tuple  $t$
- Similarly,  $t[A_u, A_v, \dots, A_w]$  refers to the subtuple of  $t$  containing the values of attributes  $A_u, A_v, \dots, A_w$ , respectively in  $t$
- All values are considered *atomic* (indivisible)
- Each value in a tuple must be from the domain of the attribute for that column, i.e.,  $v_i = t[A_i] \in \text{dom}(A_i)$
- Allow a special **NULL** value is used to represent values that are *unknown* or *inapplicable* to certain tuples

# Values in Tuple

- All values are considered *atomic* (indivisible):
  - Flat relational model
    - Composite and multivalued attributes not allowed
    - First normal form assumption
  - Multivalued attributes
    - Must be represented by separate relations
  - Composite attributes
    - Represented only by simple component attributes in basic relational model

# NULL Values

- **NULL** in a tuple:
  - Represent the values of attributes that may be unknown or may not apply to a tuple
  - Meanings for **NULL** values
    - Value unknown
    - Value exists but is not available
    - Attribute does not apply to this tuple (also known as value undefined)
- IMPORTANT: **NULL**  $\neq$  **NULL** *may be NULL due to different causes*

# Schema Definition in SQL

- The relation schema is

Customer(customer-name, customer-street, customer-city)

or

Customer	<u>customer-name</u>	customer-street	customer-city
----------	----------------------	-----------------	---------------

- The primary key is underlined in the above

```
CREATE TABLE Customer
(
    customer-name    CHAR(20)    NOT NULL,
    customer-street  CHAR(30) ,
    customer-city    CHAR(30) ,
    PRIMARY KEY (customer-name)
)
```

# Schema Definition in SQL

- To remove a relation from an SQL database, we use the **drop table** command:

**drop table r**

- We use the alter table command to add or delete attributes to an existing relation
- All records in the relation are assigned null for a new attribute.

**alter table** customer **add** phone char(10)

**alter table** customer **drop** phone

# Summary

Informal Terms	Formal Terms
Table	Relation
Column Header	Attribute
All possible <u>Column Values</u>	Domain
Row	Tuple
Table Definition	Schema of a Relation
Populated Table	State of the Relation



# Outline

- **Relational Model**
  - Relational Model Concepts
  - **Relational Database Schemas**
- ER-to-Relational Mapping
  - Translating traditional ER diagrams
  - Translating Class Hierarchy

# Relational Databases

- **Relational database schema  $S$**

- Set of relation schemas  $S = \{R_1, R_2, \dots, R_m\}$
- Set of integrity constraints (ICs)

- **Relational database state**

- Set of relation states  $DB = \{r_1, r_2, \dots, r_m\}$   $r \subseteq \text{dom}(A_1) \times \dots \times \text{dom}(A_n)$
- Each  $r_i$  is a state of  $R_i$  and such that the  $r_i$  relation states satisfy integrity constraints specified in IC
- **Invalid state**: Does not obey all the integrity constraints ICs
- **Valid state**: Satisfies all the constraints in the defined set of integrity constraints ICs

# Relational Integrity Constraints (IC)

- Constraints are **conditions** that must hold on **all** valid relation states.
- There are three *main types* of constraints in the relational model:
  - **Domain** constraint
    - Every value in a tuple must be from the domain of its attribute  
(or it could be **NULL**, if allowed for that attribute)
  - **Key** constraints
  - **Entity integrity** constraints
  - **Referential integrity** constraints ?
- ICs are specified when the schema is defined
- ICs are checked when relations are modified

# Domain Constraints

- The value of an attribute is limited to its domain.
- A domain can impose rules on both formats and valid value ranges.
  - A salary value cannot be negative
  - An employee's name cannot be **NULL**
    - This is called the **NOT NULL** constraint

- Example:

Constraint name (optional): if the constraint is violated, the constraint name is returned and can be used to identify the error

```
CREATE DOMAIN hourly-wage NUMERIC(5,2)
CONSTRAINT wage-value-test CHECK (value > 4.00)
```

- The domain hourly-wage is a decimal number with 5 digits, 2 of which are placed after the decimal point
- The domain has a constraint that ensures that the hourly wage is greater than 4.00

# Key Constraints

- Certain minimal subset of the fields (candidate key) of a relation is a unique identifier for a record
- Out of all the available candidate keys, a database designer can identify a primary key
  - The DBMS may create an index with the primary key
- **Key of R:**
  - Is a set of attributes K of R with the following condition:
    - No two tuples in any valid relation state  $r(R)$  will have the same value for K (that is, for any distinct tuples  $t_1$  and  $t_2$  in  $r(R)$ ,  $t_1[K] \neq t_2[K]$ ) *no duplicate key tuples*
- **Candidate Key of R:**
  - A "minimal" key

# Key Constraints

- If a relation has several **candidate keys**, one is chosen to be the **primary key**
  - The primary key attributes are underlined
- The primary key is used to *reference* the tuple from another tuple
  - General rule:
    - choose as primary key the smallest of the candidate keys (in terms of space) *space*
    - choice is sometimes subjective

# Key in SQL

- In SQL, we can declare
  - a key by the **UNIQUE** command
  - a primary key by the **PRIMARY KEY** constraint

```
CREATE TABLE Students  
(
```

```
    sid    CHAR(20) ,  
    name   CHAR(30) ,  
    login  CHAR(20) ,  
    age     INTEGER ,  
    gpa     REAL ,
```

```
    UNIQUE (name, age) ,
```

```
    CONSTRAINT StudentsKey PRIMARY KEY (sid)
```

```
)
```

Primary key

key

Constraint name (optional)

# Entity Integrity

*minimal set of key  $\Rightarrow$  CK (single attribute)  $\Rightarrow$  one of them PK (single attribute)*

- The *primary key attributes* PK of each relation schema R


in S **CANNOT** have **NULL** values in any tuple of  $r(R)$

- This is because primary key values are used to identify the individual tuples
- $t[PK] \neq \mathbf{NULL}$  for any tuple  $t$  in  $r(R)$
- If PK has several attributes, **NULL** is not allowed in any of these attributes

**No primary key value can be NULL**



# Referential Integrity

- Referential Integrity (RR) constraint specified between two relations. ~~M~~aintains consistency among tuples in two relations
- Tuples in the **referencing relation** R1 have attributes FK (called **foreign key** attributes) that reference the **primary key** attributes PK of the **referenced relation** R2
  - A tuple t1 in R1 is said to **reference** a tuple t2 in R2 if
$$t1[FK] = t2[PK]$$
- A RR constraint can be displayed in a relational database schema as a directed arc  **FROM** R1.FK **TO** R2.PK.

# Foreign Key Constraint

- A **foreign key** is a set of attributes in one relation  $R1$  that is used to refer to a tuple in another relation  $R2$
- Statement of the constraint
  - The value in FK of the **referencing relation**  $R1$  can be **either**:
    - (1) a value of an existing primary key value of a corresponding primary key PK in the **referenced relation**  $R2$ , or
    - (2) a **NULL**
  - In case (2), the FK in  $R1$  should **NOT** be a part of its own primary key (because of Entity Integrity Constraint)

# Example 1: Relational Schema Diagram

- Student(Student-id, Student-Name)  $\neq \text{NULL}$  no foreign keys
- Take(Student-id, Course-id, semesterNo) PK: (stu-id, course-id)
- Course(Course-id, Course-Name)  $\neq \text{NULL}$  no foreign keys

(Student-id, Course-id) in relation Take is a primary key

Student-id, course-id in relation Take are foreign keys

Student 

stu-id	stu-name
--------	----------

 $\neq \text{NULL}$  no foreign keys

foreign key<sub>1</sub> = Take.stu-id  $\rightarrow$  Student.stu-id

Draw a relational schema diagram specifying the foreign

Take 

stu-id	course-id	semester-no
--------	-----------	-------------

keys for this schema

Course

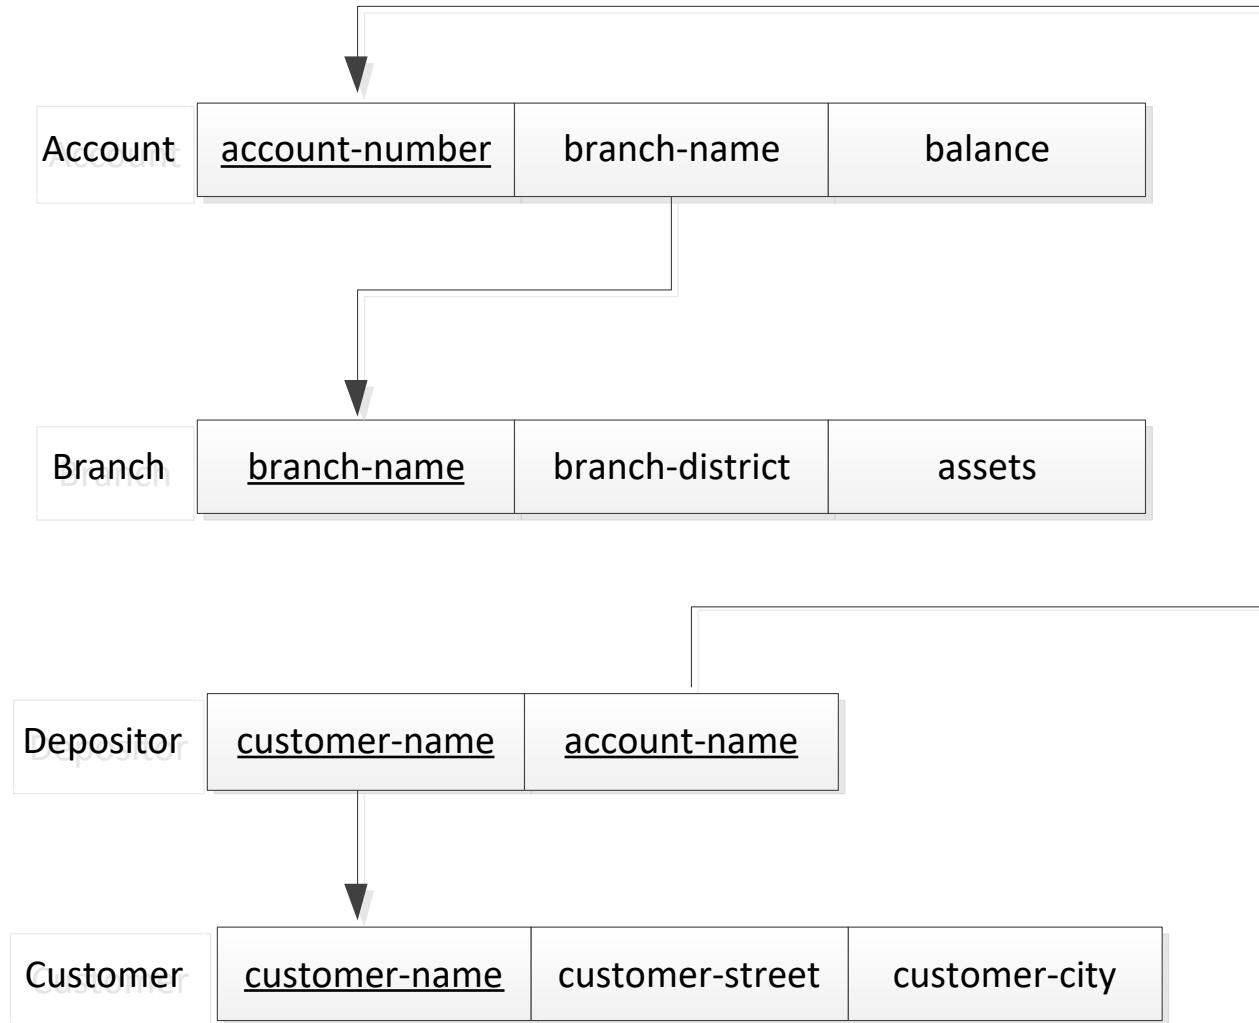
course-id	course-name
-----------	-------------

 $\neq \text{NULL}$

no foreign keys

FK<sub>2</sub> = Take.course-id  $\rightarrow$  course.course-id

## Example 2: Foreign Key



## Example 2: Foreign Key

Account

<u>account-number</u>	branch-name	balance
-----------------------	-------------	---------

```
CREATE TABLE account
```

```
(
```

```
  account-number    CHAR(10)    NOT NULL,
```

```
  branch-name       CHAR(15) ,
```

```
  balance           INTEGER,
```

```
  PRIMARY KEY (account-number)
```

```
  FOREIGN KEY (branch-name) REFERENCES branch
```

```
)
```

Branch

<u>branch-name</u>	branch-district	assets
--------------------	-----------------	--------

## Example 2: Foreign Key

Depositor

customer-name

account-number

```
CREATE TABLE depositor
```

```
(
```

```
customer-name      CHAR(20)  NOT NULL,
```

```
account-number     CHAR(10)  NOT NULL,
```

```
PRIMARY KEY (customer-name, account-number),
```

```
FOREIGN KEY (customer-name) REFERENCES customer
```

```
FOREIGN KEY (account-number) REFERENCES account
```

```
)
```

Account

account-number

branch-name

balance

Customer

customer-name

customer-street

customer-city

# Populated Database State

- Each *relation* will have many tuples in its current relation state
- The **relational database state** is a union of all the individual relation states
- Whenever the database is changed, a new state arises
- Basic operations for changing the database:
  - INSERT a new tuple in a relation
  - DELETE an existing tuple from a relation
  - MODIFY an attribute of an existing tuple

# Possible Violations for INSERT

- Domain constraint:
  - if one of the attribute values provided for the new tuple is not of the specified attribute domain
- Key constraint:
  - if the value of a key attribute in the new tuple already exists in another tuple in the relation
- Referential integrity:
  - if a foreign key value in the new tuple references a primary key value that does not exist in the referenced relation
- Entity integrity:
  - if the primary key value is null in the new tuple



# Example:

Account(account-number, branch-name, balance)

Branch(branch-name, branch-district, assets)

Depositor(customer-name, account-number)

Customer(customer-name, customer-street, customer-city)

cannot add ( KLN, 111, 3 ) to Account

cannot add ( Chris, 222 ) to Depositor

cannot add ( Mary, 999 ) to Depositor

Customer

customer-name	customer-street
Tim	Main
Smith	North
Mary	North

Branch

branch-name	branch-district	assets
CUHK	Shatin	1000
CMTR	Central	2000
SKCR	Shatin	4000

Account

branch-name	account-number	balance
CUHK	222	2
CUHK	777	5
CMTR	333	1
SKCR	444	5
SKCR	888	5

Depositor

customer-name	account-number
Tim	222
Mary	888
Tim	444
Smith	777
Tim	333

# Possible Violations for DELETE

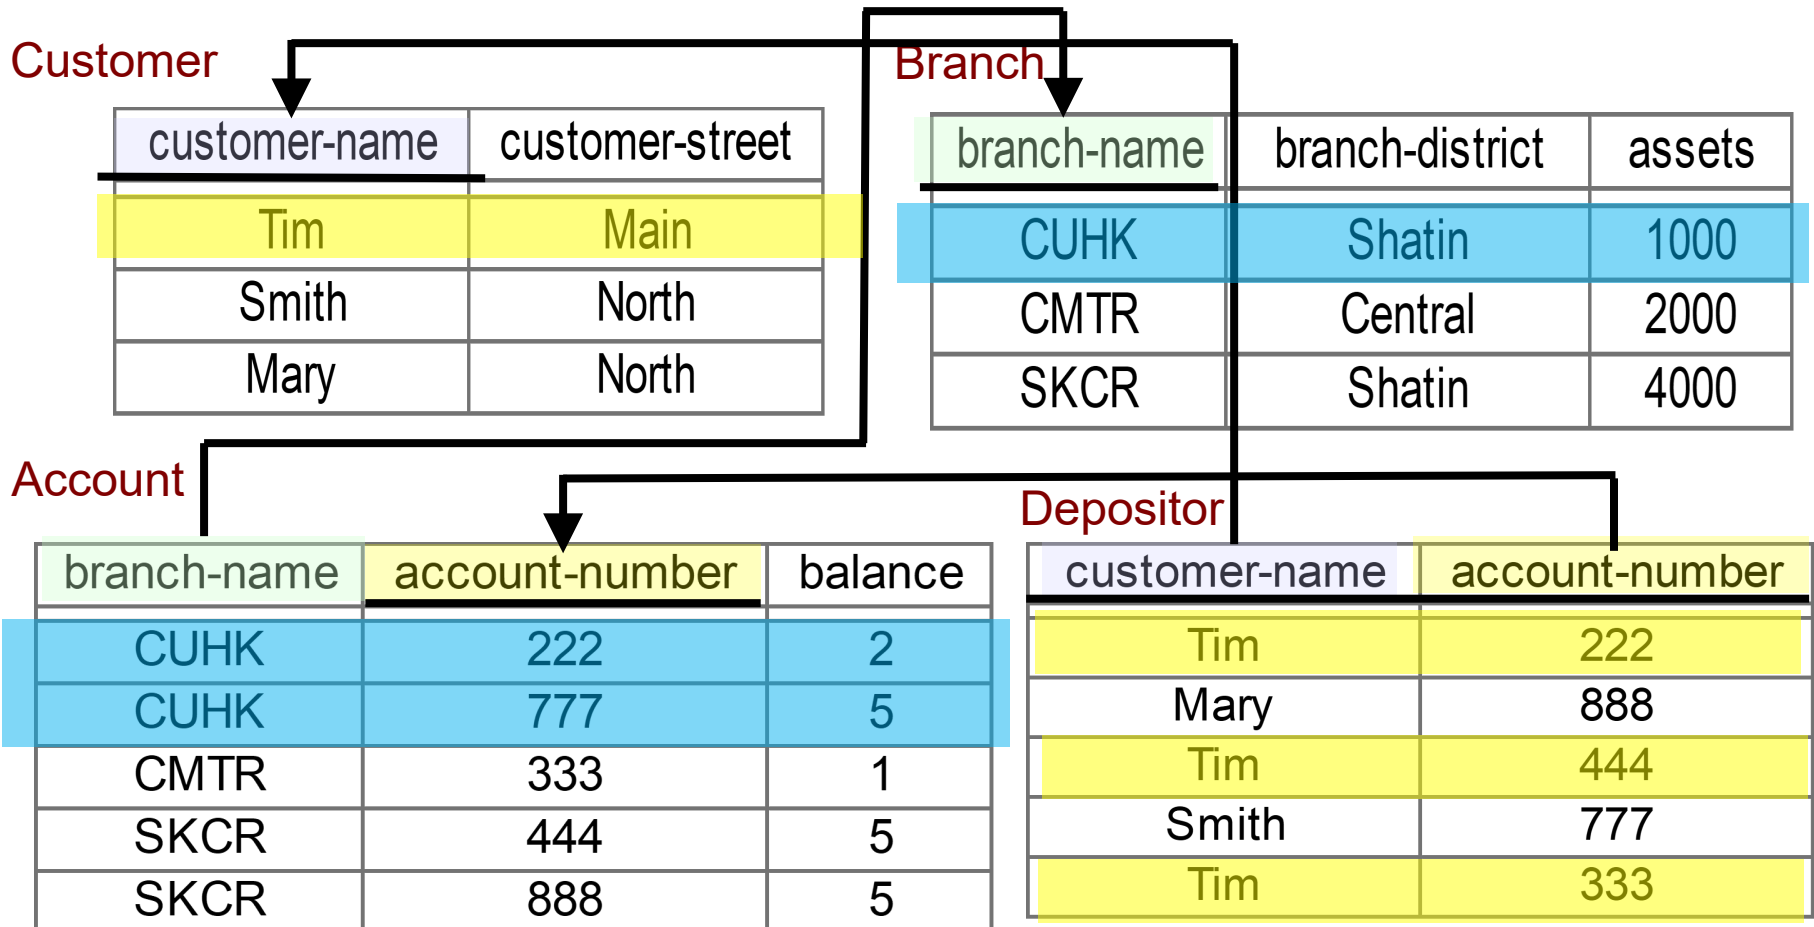
- Referential constraints
  - If the primary key value of the tuple being deleted is referenced from other tuples in the database
- Options of remedies:
  - **RESTRICT**: reject the deletion
  - **CASCADE**: delete the record in the referencing table
  - **SET NULL**: set the foreign keys of the referencing tuples to NULL
  - **SET DEFAULT**: set the foreign keys of the referencing tuples to default value

# Example:

Account(account-number, branch-name, balance)  
Branch(branch-name, branch-district, assets)  
Depositor(customer-name, account-number)  
Customer(customer-name, customer-street, customer-city)

Cannot simply delete Customer Tim

Cannot simply delete branch CUHK



# Example: Foreign Key

- Cascading delete:
  - Delete Customer Tim, cascaded delete on depositor
  - Delete branch CUHK, cascaded delete on account

```
CREATE TABLE account
(
    branch-name          CHAR(15) ,
    account-number       CHAR(10)   NOT NULL ,
    balance              INTEGER ,
    PRIMARY KEY (account-number) ,
    FOREIGN KEY (branch-name) REFERENES branch
                        ON DELETE CASCADE
                        ON UPDATE CASCADE
                        [ON UPDATE NO ACTION]
)
```

# Example: Foreign Key

Branch

branch-name

branch-district

assets

```
CREATE TABLE branch (  
    branch-name      CHAR(20)      DEFAULT 'CUHK' ,  
    branch-district CHAR(30) ,  
    assets           INTEGER ,  
    PRIMARY KEY (branch-name) )
```

```
CREATE TABLE account (  
    branch-name      CHAR(15)      DEFAULT 'CUHK' ,  
    account-number   CHAR(10)      NOT NULL ,  
    balance          INTEGER ,  
    PRIMARY KEY (account-number) ,  
    FOREIGN KEY (branch-name) REFERENCES branch  
        ON DELETE CASCADE  
        ON UPDATE CASCADE )  
    [ON DELETE SET DEFAULT]  
    [ON DELETE SET NULL]
```

# Possible Violations for UPDATE

- Constraint violations depending on the attribute being updated:
  - Updating the primary key (PK):
    - Similar to a DELETE followed by an INSERT
    - Need to specify similar options to DELETE
  - Updating a foreign key (FK):
    - May violate referential integrity
  - Updating an ordinary attribute (neither PK nor FK):
    - Can only violate domain and business rules constraints

# Outline

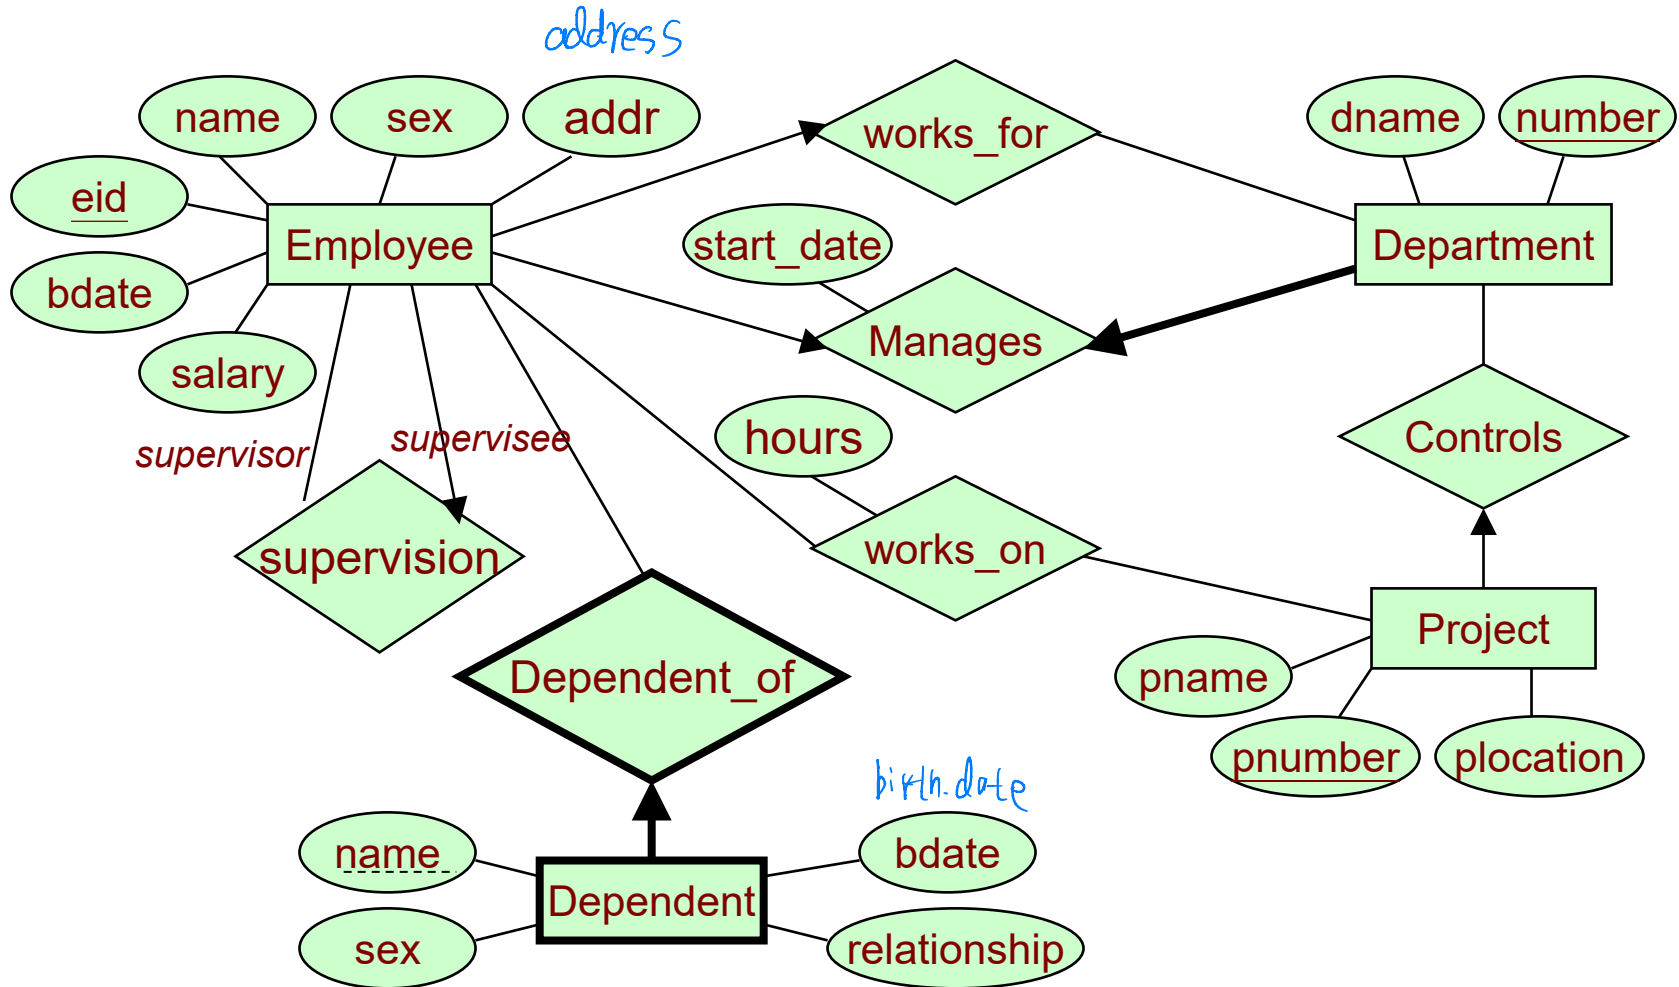
- Relational Model
  - Relational Model Concepts
  - Relational Database Schemas
- **ER-to-Relational Mapping**
  - **Translating traditional ER diagrams**
  - Translating Class Hierarchy

# ER-to-Relational Mapping

- Typically, database designers begin with the ER model, which is very **expressive** and **user-friendly** to human
- Then, the ER model is mapped to the relational model for DBMS manipulations
- Database queries and updates will be written according to the relational model



# Running Example



# ER to Relational Model

## 1. Convert entities first

- From Strong Entity to Weak Entity

## 2. Convert relations

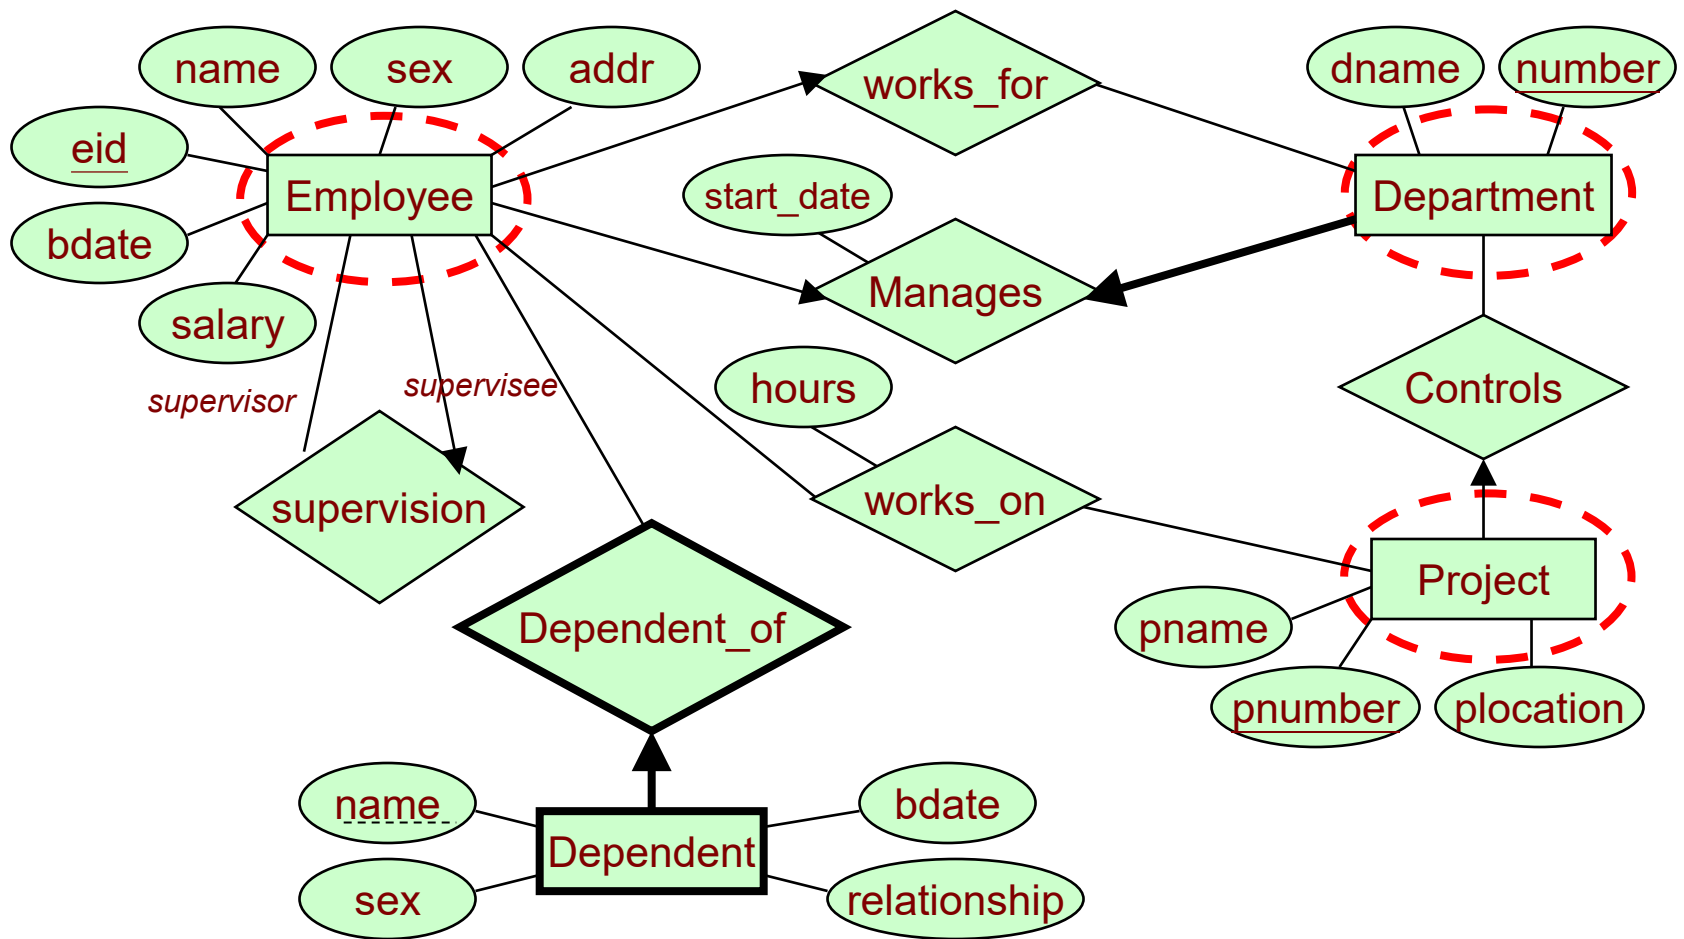
- From 1-to-1, 1-to-many (many-to-1) to many-to-many
- From binary relation to non-binary relation




# Steps

- ER-to-Relational Mapping Algorithm
  - **Step 1:** Mapping of Regular Entity Types
  - **Step 2:** Mapping of Weak Entity Types
  - **Step 3:** Mapping of Binary 1:1 Relation Types
  - **Step 4:** Mapping of Binary 1:N Relationship Types.
  - **Step 5:** Mapping of Binary M:N Relationship Types.
  - **Step 6:** Mapping of N-ary Relationship Types.

# Step 1 (Strong Entity)

- For each strong entity set  $E$  in the ER schema,
  - create a relation schema  $R$  that includes all the attributes of  $E$
  - choose one set of key attributes of  $E$  as a primary key for  $R$
  - if the chosen key of  $E$  is composite, the set of simple attributes that form it will together form the primary key of  $R$





# Step 1 (Strong Entity)

- Example
- We create the relation schemas EMPLOYEE, DEPARTMENT and PROJECT

EMPLOYEE	name	<u>eid</u>	bdate	addr	sex	salary
----------	------	------------	-------	------	-----	--------


DEPARTMENT	dname	<u>dnumber</u>
------------	-------	----------------

PROJECT	pname	<u>pnumber</u>	plocation
---------	-------	----------------	-----------

# Step 1 (Strong Entity)

*(omitted)*

- If there is a derived attribute, what should we do?



We have two choices.

Choice 1: Include this derived attribute

Adv: We can directly obtain the value of the derived attribute

Disadv: We may encounter some data inconsistencies

~~Choice 2:~~ NOT include this derived attribute


Adv: We can avoid data inconsistency

Disadv: We need to perform some operations to obtain the value of the derived attribute



# Step 1 (Strong Entity)

- If there is a composite attribute, what should we do?



We have two choices.


Choice 1: Include the high-level attribute only (e.g., address)

Choice 2: Include all low-level attributes (e.g., street, city, country)

*prefer*

# Step 1 (Strong Entity)

- If there is a multi-valued attribute, what should we do?



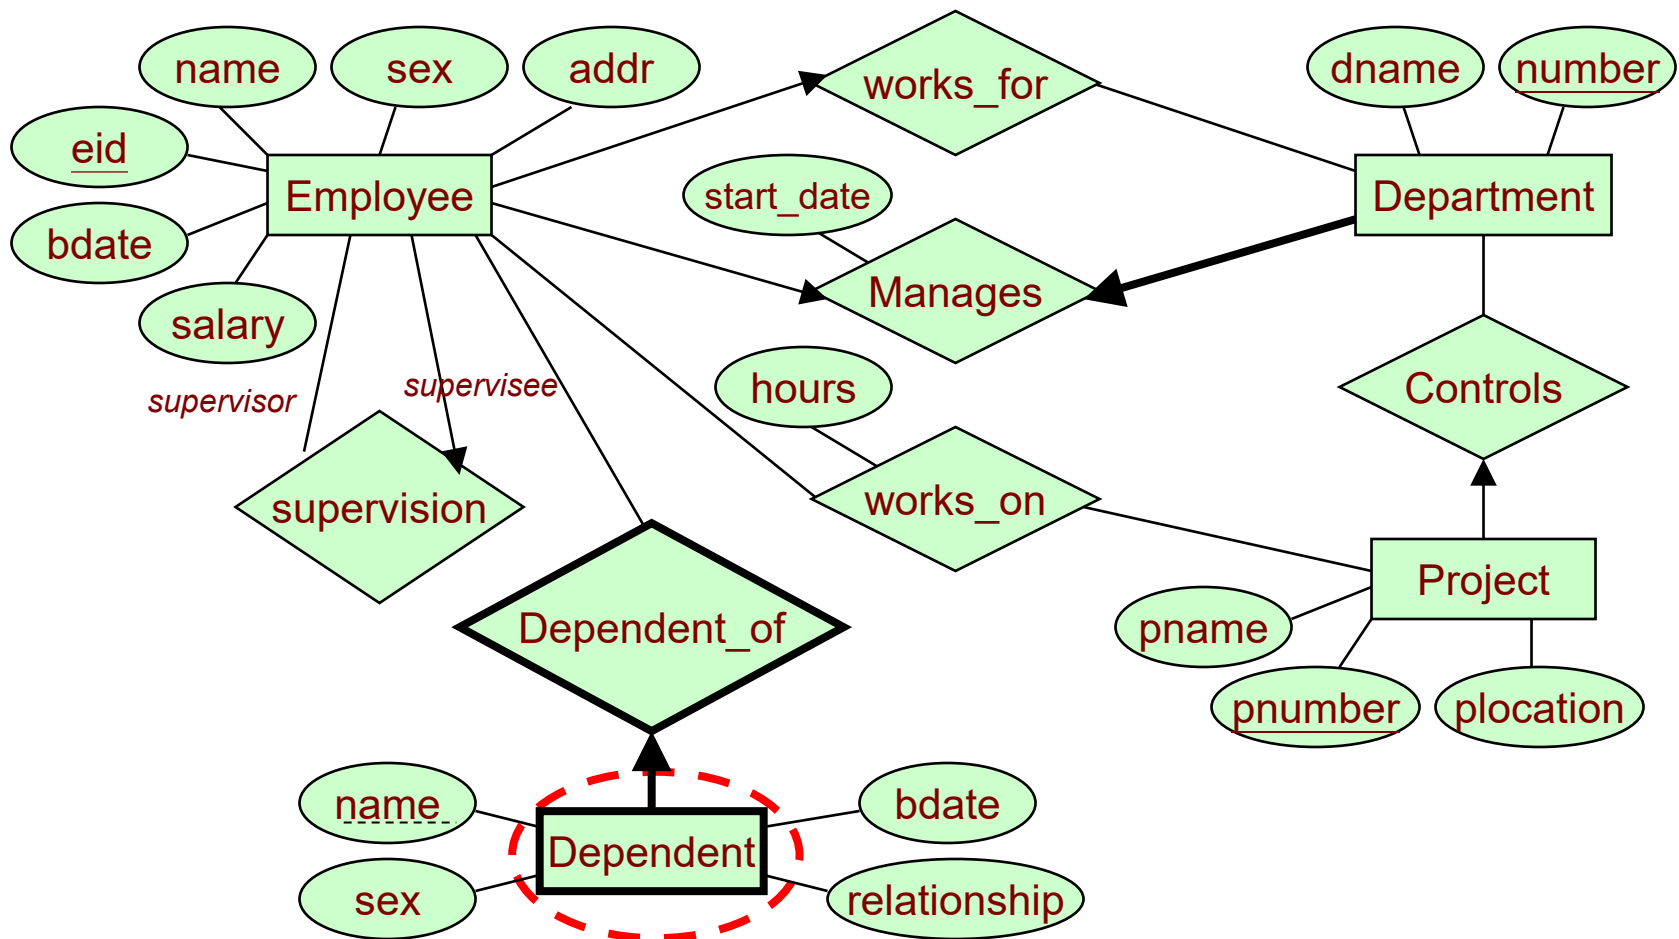
We have two choices.

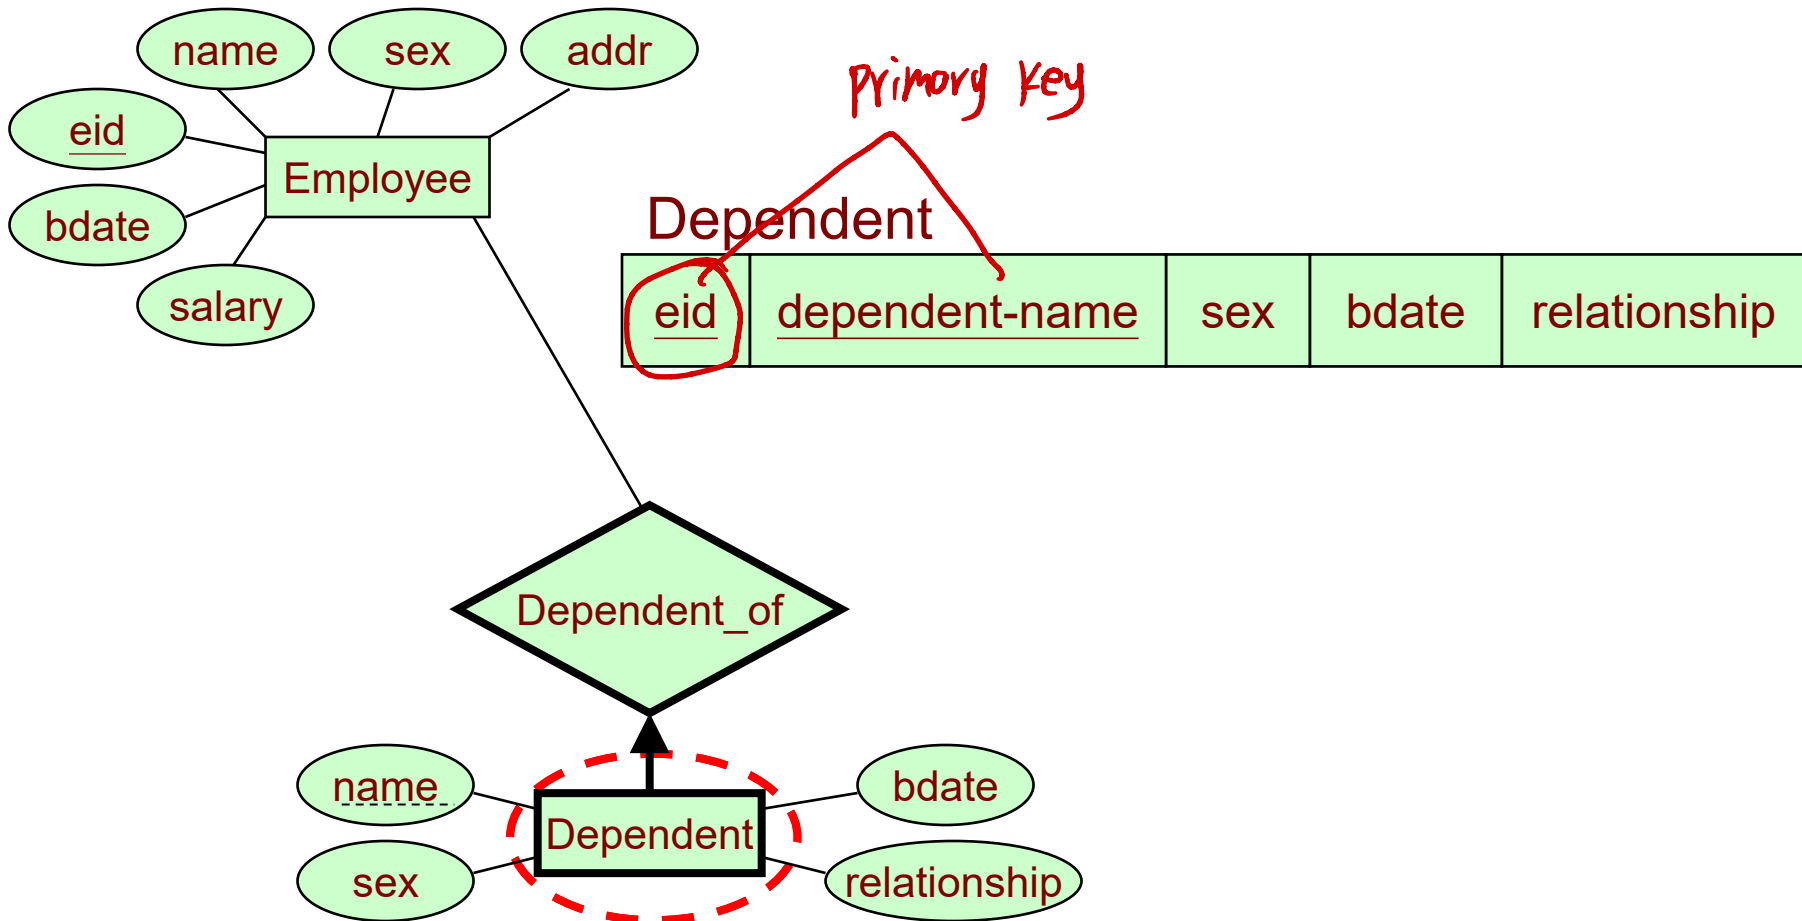
Choice 1: Include one attribute only (e.g., phone)

✓ Choice 2: Create another table containing the primary key of the entity set and the multi-valued attribute. e.g., create a schema PhoneTable (eid, phone) *foreign key*

## Step 2 (Weak Entity)


- For each **weak entity set** W in the ER model,
  - Create a relation schema R, and include all attributes
  - In addition, include the primary key(s) of the owner(s)
  - The primary key of R is the combination of the primary key(s) of the owner(s) and the discriminator  
partial key of the weak entity set W



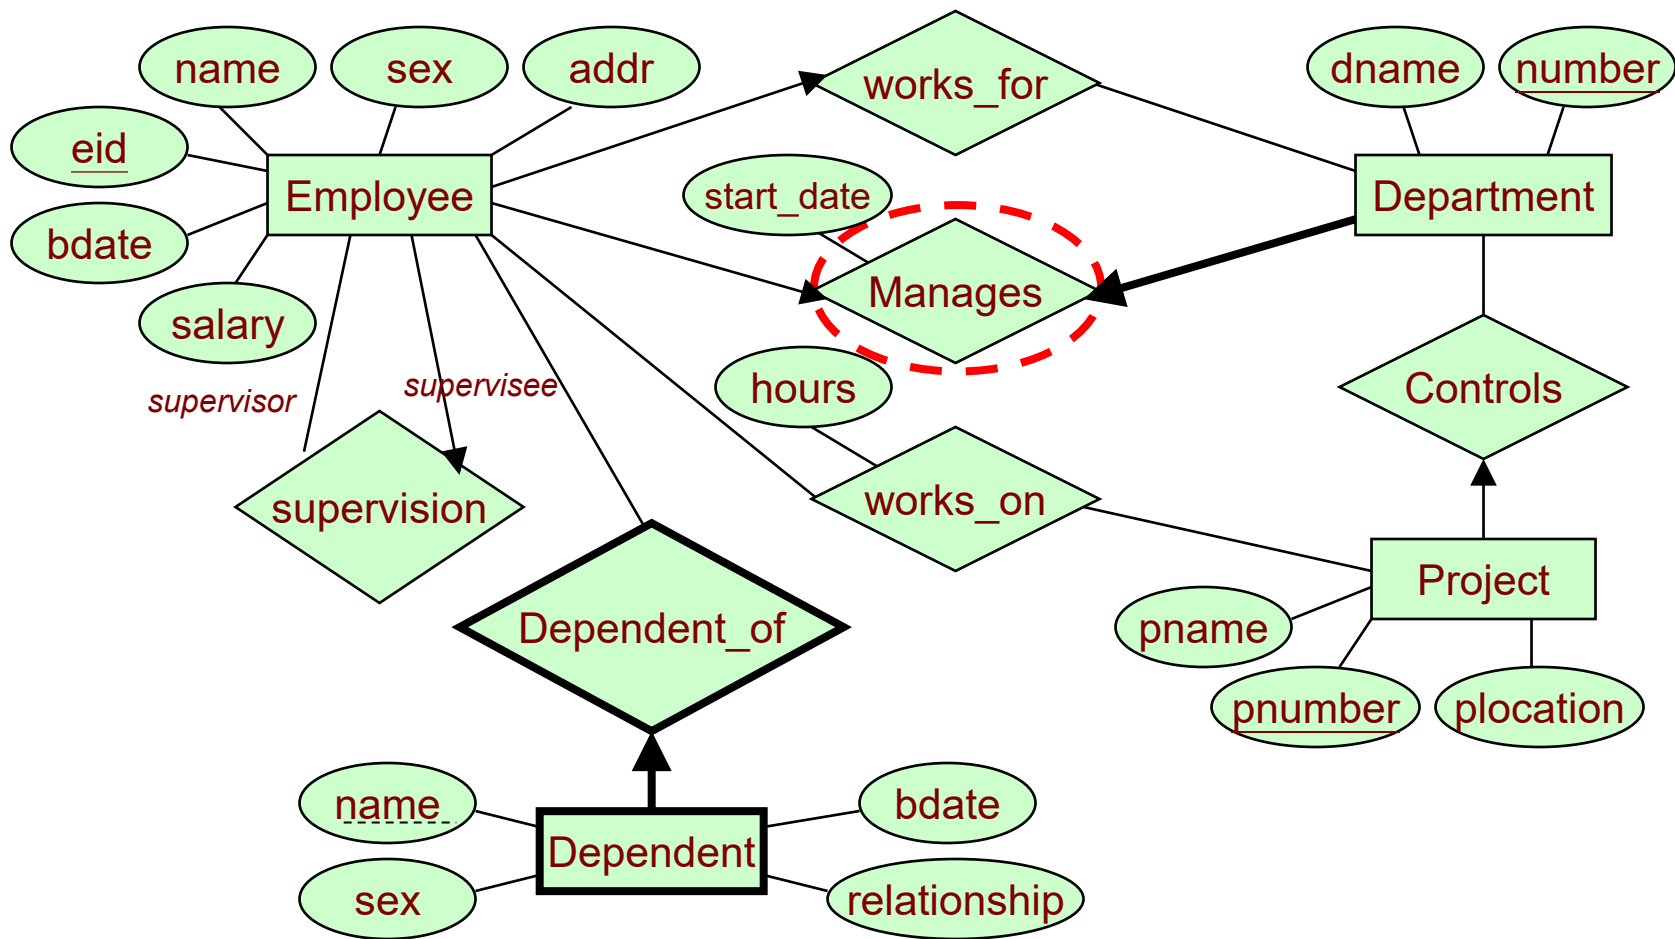


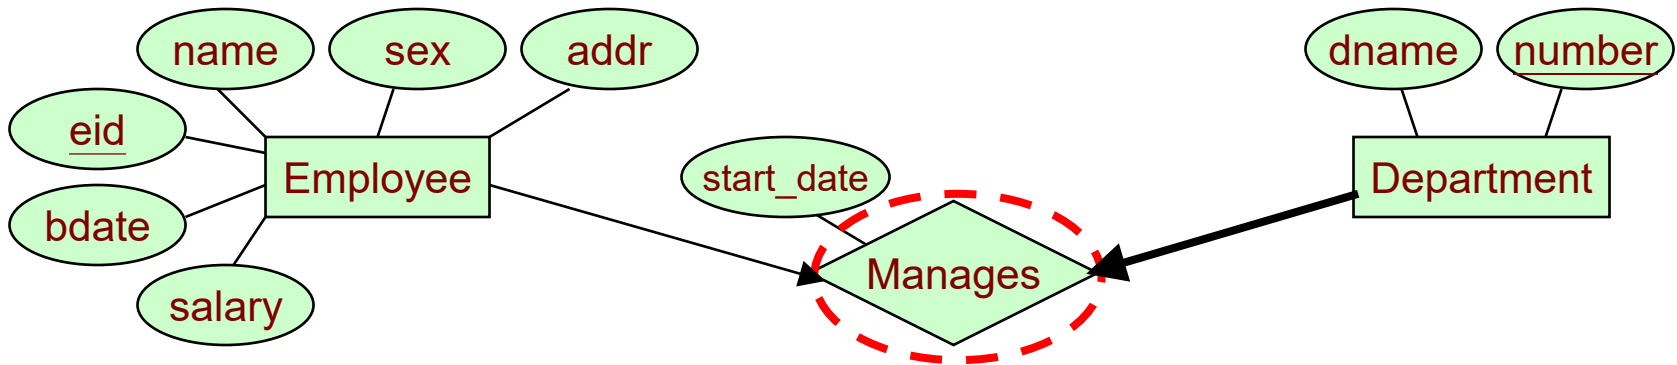
# Step 3 (1-to-1 Relationship)

- For each binary one-to-one (1:1) relationship set R



- Choose one of the 2 relation schemas, say S,
  - get primary key of T,
  - include it as foreign keys in S
- Better /if S has total participation in R
- Include the attributes of the relationship set R as attributes of S






- We include the **primary key** of EMPLOYEE as foreign key in DEPARTMENT and rename it **mgr\_id**
- We include the attribute **startdate** of MANAGES and rename it **mgr\_start\_date**

DEPARTMENT

dname	<u>dnumber</u>	mgr_id	mgr_start_date
-------	----------------	--------	----------------



Compare the following two choices to include MANAGES:



Add information to EMPLOYEE

name	<u>id</u>	bdate	addr	sex	salary	dnum	sdate
Yeung	7	080370	...	F	20K	3	010100
Chan	3	031060	...	M	30K	4	020399
Wong	4	010280	...	F	10K	Null	Null
Cheung	8	220985	...	M	24K	Null	Null

dname	<u>dnumber</u>	mgr_id	mgr_start_date
Personnel	4	3	020399
Marketing	3	7	010100
...	...	...	...

Add to DEPARTMENT

- In the above, the NULL value is a special value meaning that the value is either unknown or not applicable
- Notice that an alternative mapping of a one-to-one relationship set is possible by merging the two entity sets and the relationship into a single relation
- This is appropriate when both participations are total

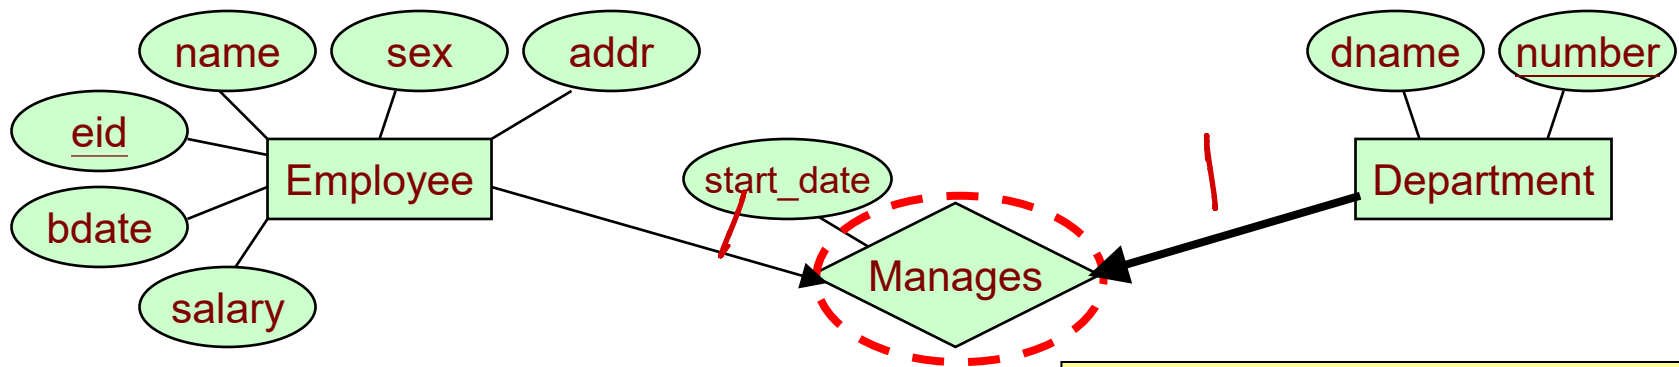
Mapping 1-to-1 Relationship:

Advantage:

The total number of relations remain unchanged

Disadvantage:

It may store NULL values if there is no total participation



Can we create a new relation

Manages (eid, number, start\_date)

or

Manages (eid, number, start\_date)

for this relationship?

Yes.

It can be used if there are only  
a few relationship instances

Advantage:

It can avoid storing NULL values  
if there is no total participation


Disadvantage:

There is one additional  
relation

- There are three approaches for mapping 1-1 binary relationship:
  1. **Foreign Key approach:** Choose one of the relations-say S-and include a foreign key in S the primary key of T. It is better to choose an entity type with total participation in R in the role of S
  2. **Merged relation option:** An alternate mapping of a 1:1 relationship type is possible by merging the two entity types and the relationship into a single relation. This may be appropriate when both participations are total
  3. **Cross-reference or relationship relation option:** The third alternative is to set up a third relation R for the purpose of cross-referencing the primary keys of the two relations S and T representing the entity types

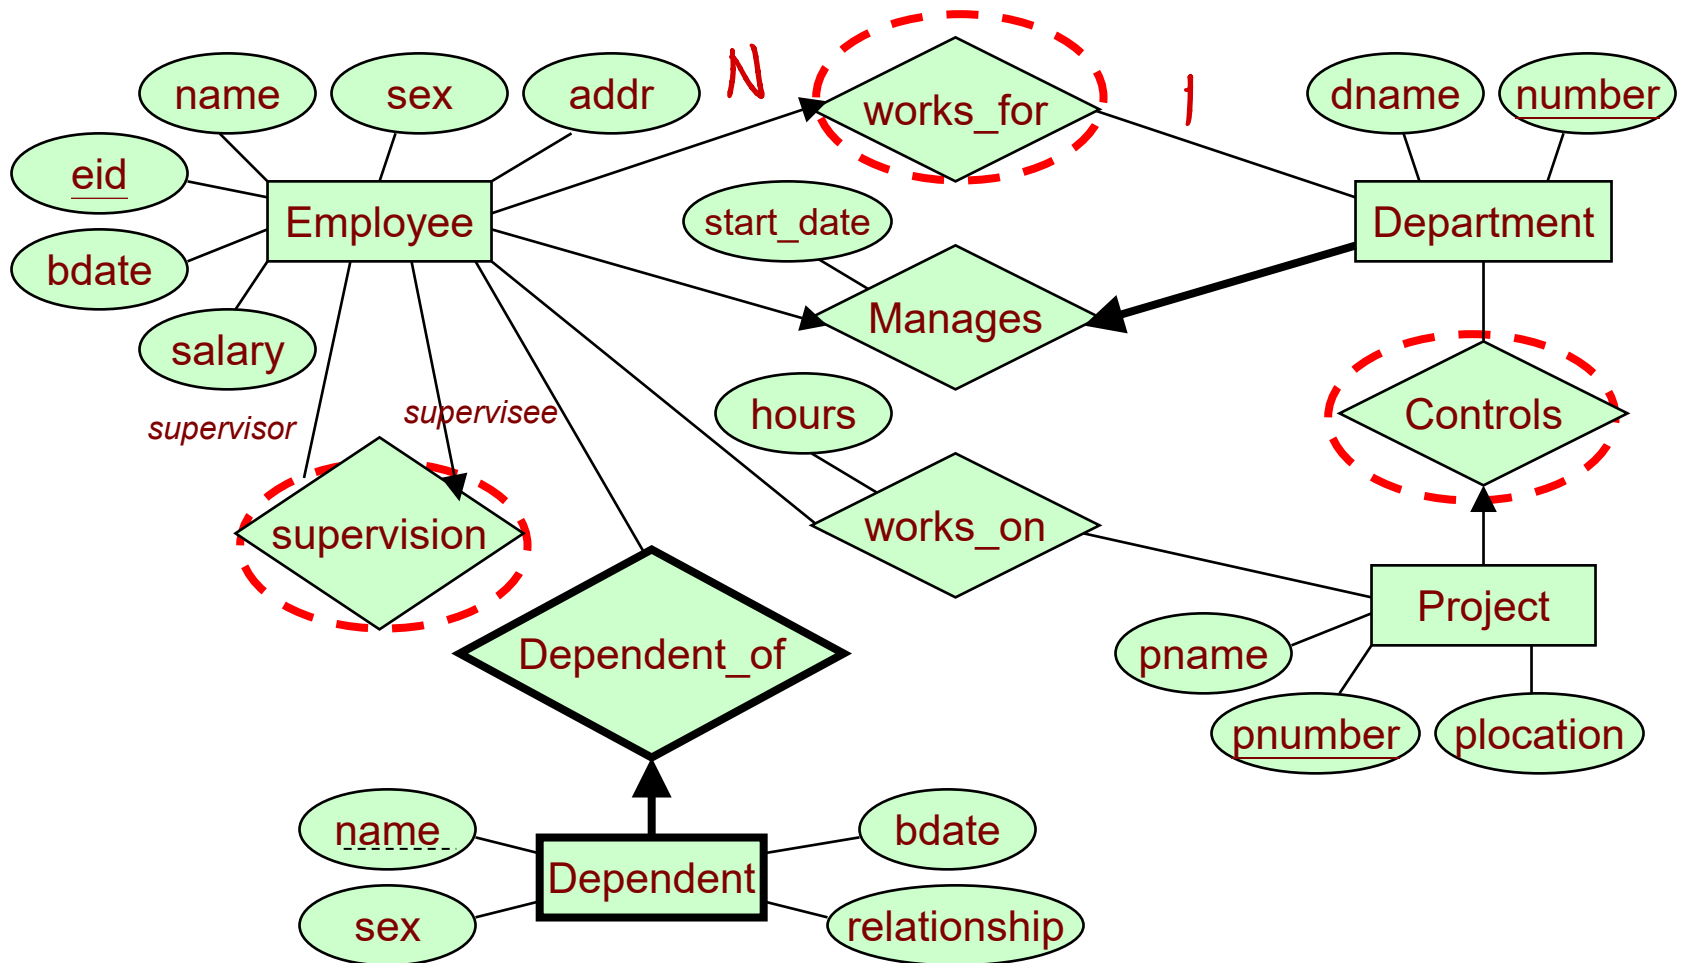
# Step 4 (1-to-many Relationship)

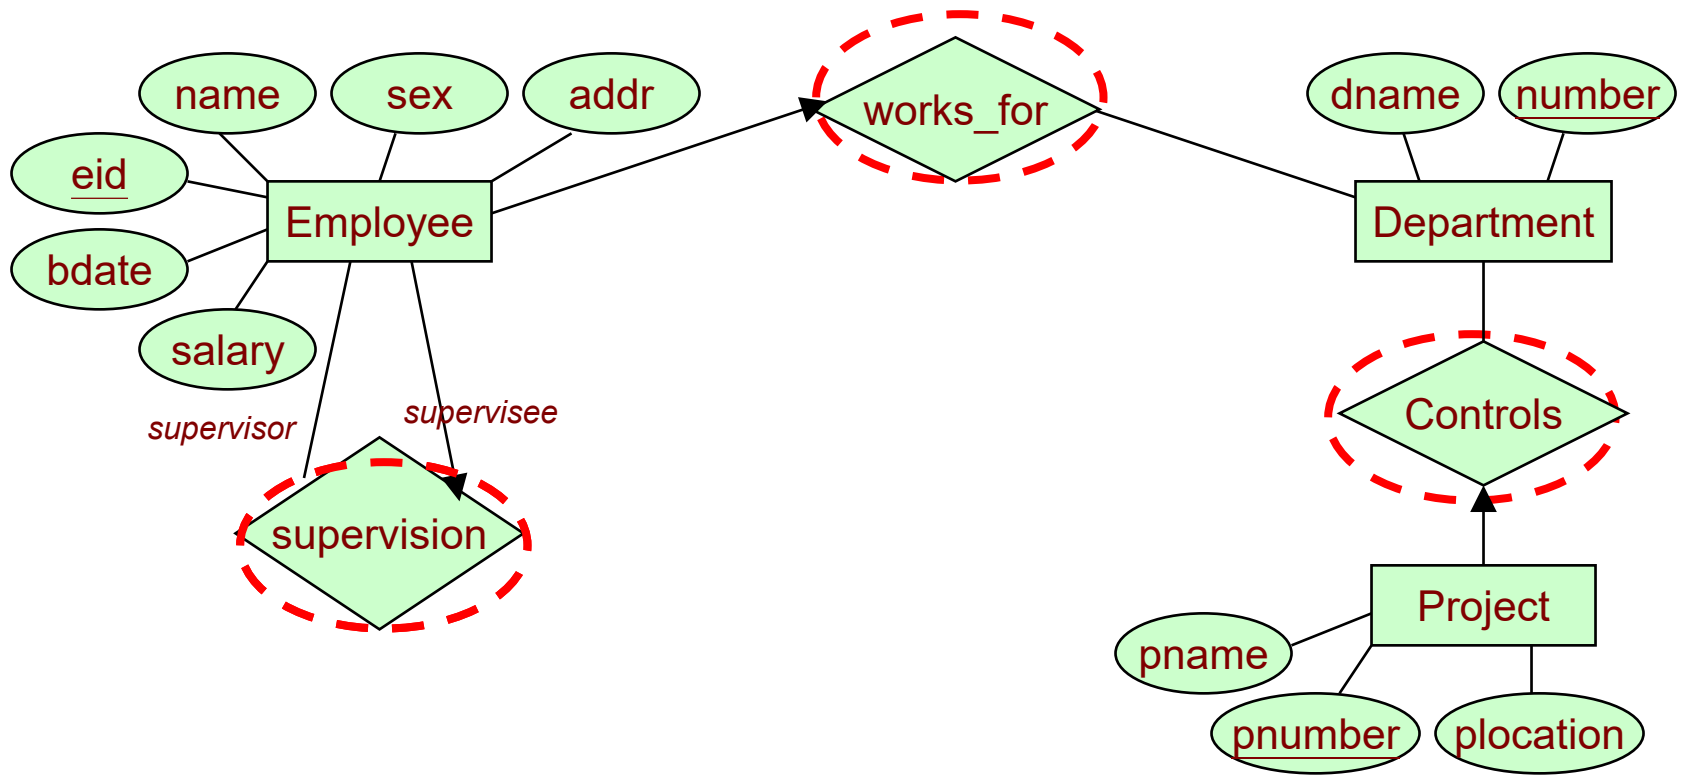
- For each **binary one-to-many** relationship set



- Include as foreign key in S the primary key that represents the other entity set T participating in R
- Include any attributes of the one-to-many relationship set as attributes of S
- In other words, we always choose the one who determines the relationship as R

Add primary key of entity(1) to the one(N) as a foreign key.






- The primary key dnumber of the DEPARTMENT relation schema is included as foreign key in the EMPLOYEE relation schema
- We rename it as dno (The renaming is not necessary but makes the name more meaningful.)

EMPLOYEE

name	<u>eid</u>	bdate	addr	sex	salary	dno
------	------------	-------	------	-----	--------	-----





Compare the following 2 choices:

Add employees to department



DEPARTMENT

dname	<u>dnumber</u>	mgr_id	mgr_start_date	<u>eid</u>
Personnel	4	3	020399	3
Personnel	4	3	020399	11
Personnel	4	3	020399	12
Personnel	4	3	020399	13
Marketing	3	7	010100	21
Marketing	3	7	010100	7
Marketing	3	7	010100	22
...	...	...	...	...

EMPLOYEE

name	<u>id</u>	bdate	addr	sex	salary	dno
Yeung	7	080370	...	F	20K	3
Chan	3	031060	...	M	30K	4
Wong	4	010280	...	F	10K	7
Cheung	8	220985	...	M	24K	1

Add department to employee



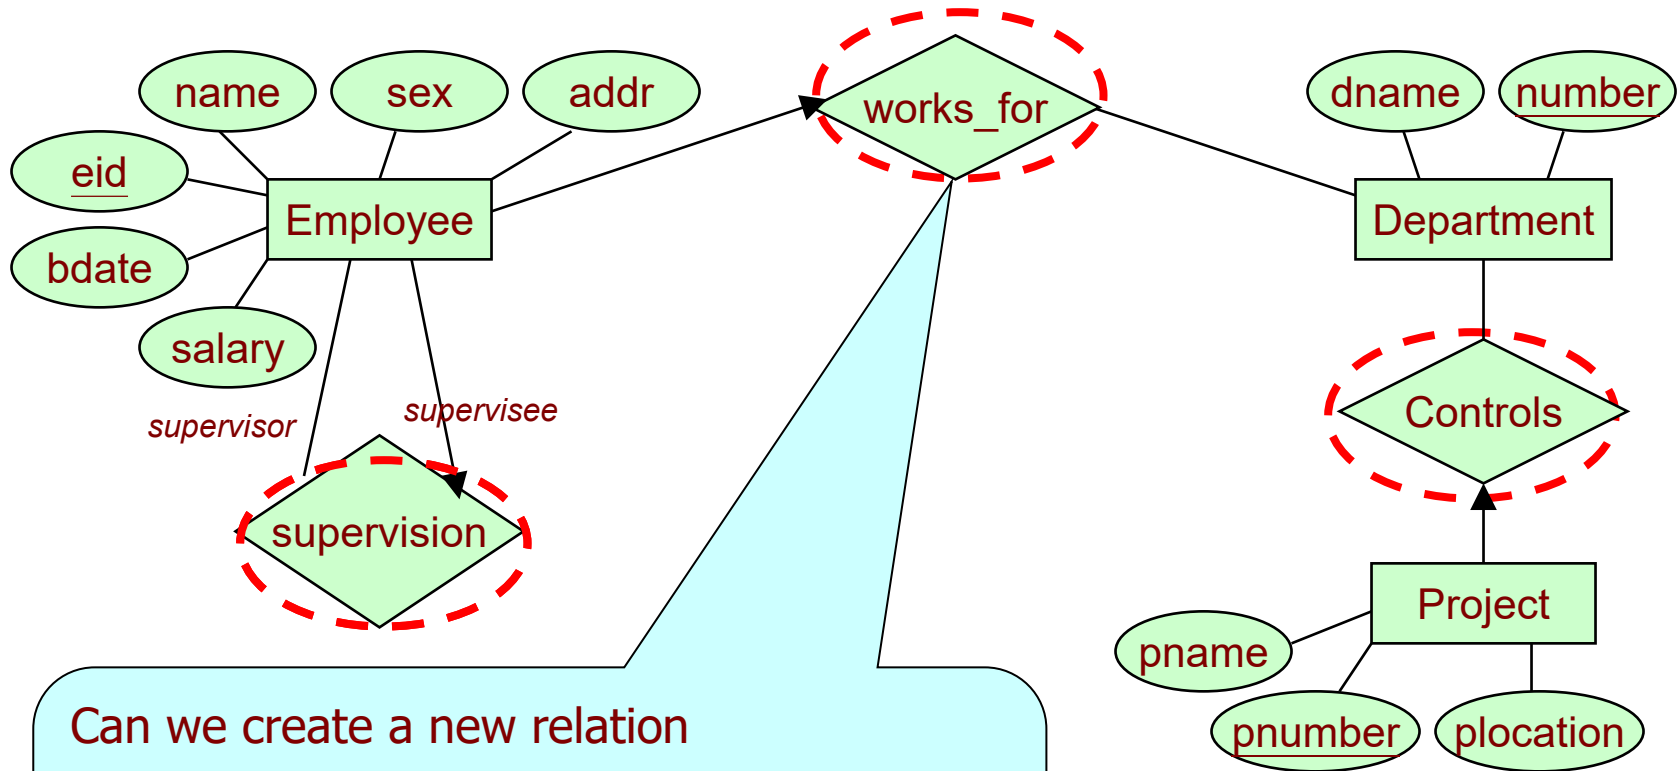
- For SUPERVISION,
  - include the primary key of the EMPLOYEE as foreign key in the EMPLOYEE, and call it super\_id

*Employee* recursive 自己添加自己

EMPLOYEE	name	<u>eid</u>	bdate	addr	sex	salary	dno	super_id
----------	------	------------	-------	------	-----	--------	-----	----------

- For CONTROLS relationship,
  - include dnum as foreign key in PROJECT,
  - which references the primary key dnumber of DEPARTMENT

PROJECT	pname	<u>pnumber</u>	plocation	dnum
---------	-------	----------------	-----------	------



Can we create a new relation


`works_for (eid, number)`

for this relationship?

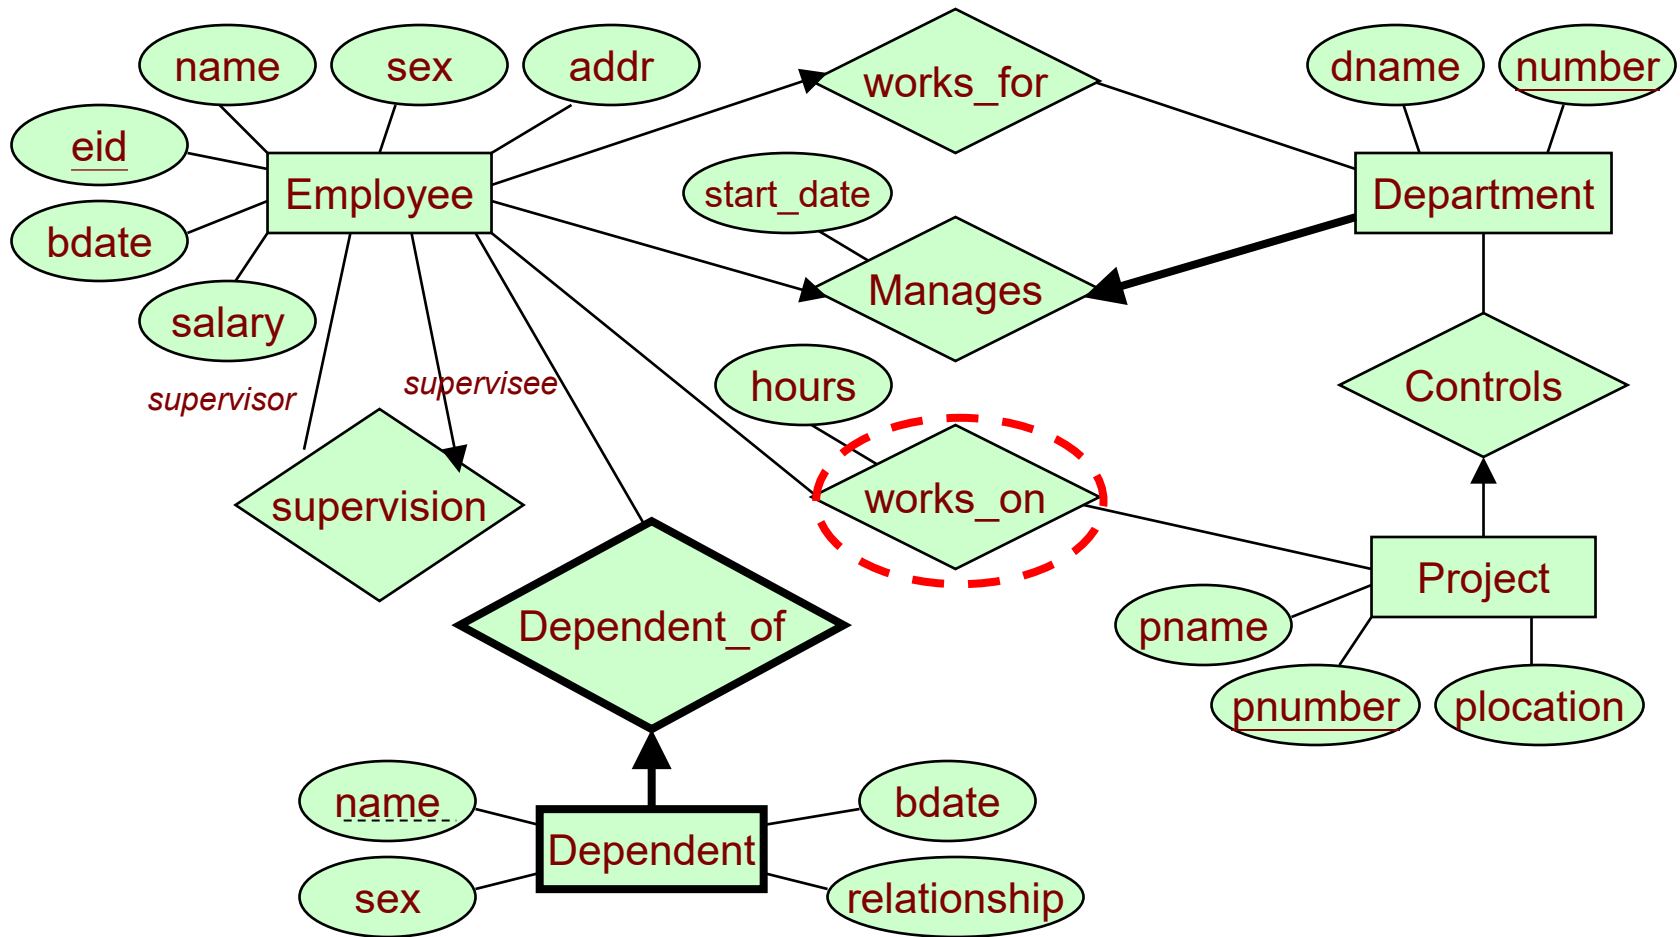
Yes.

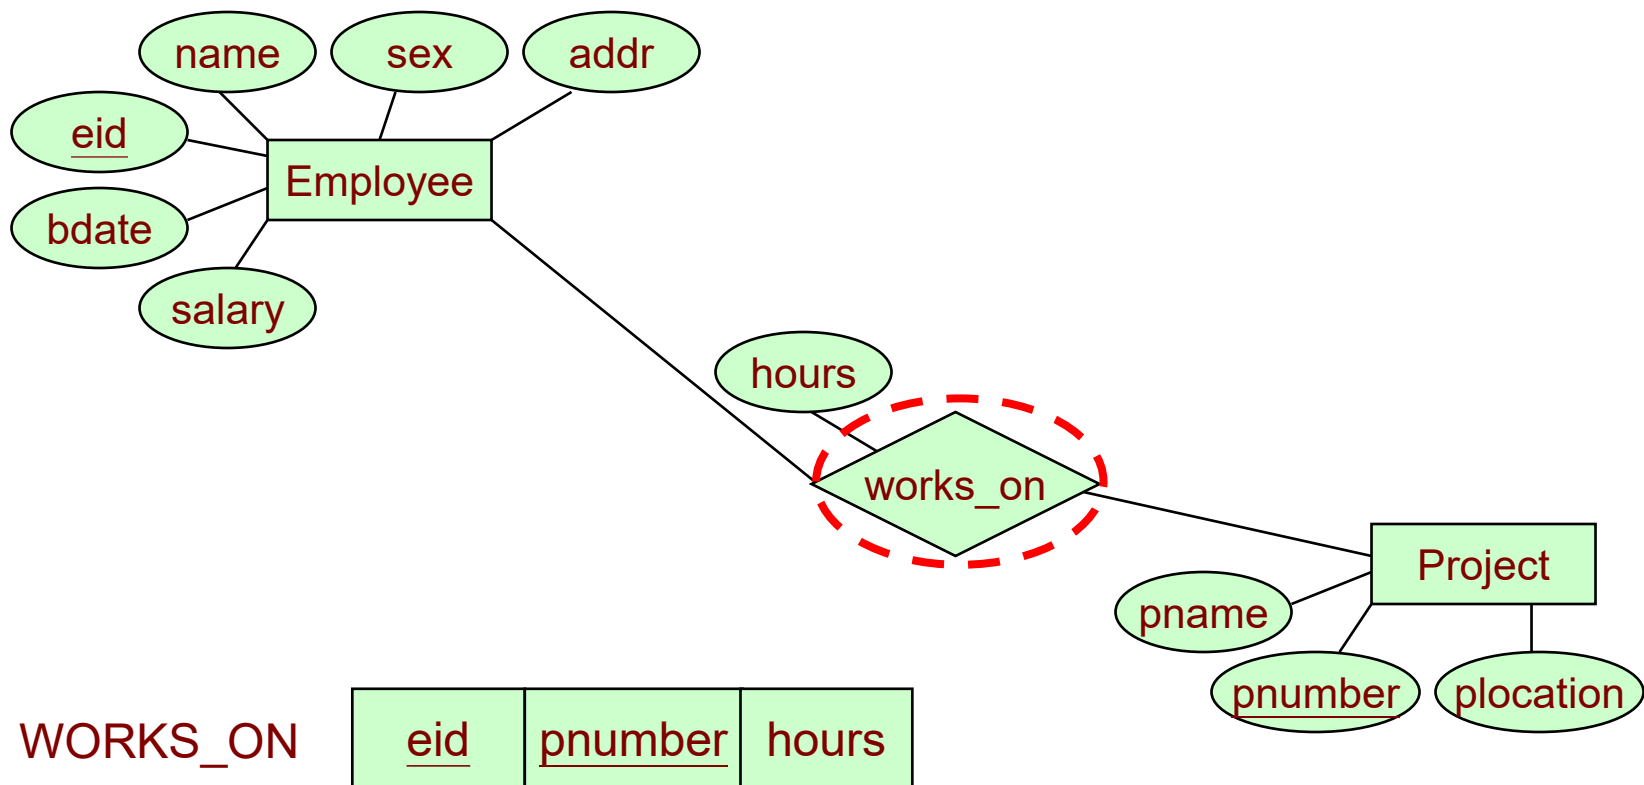
# Step 5 (Many-to-many Relationship)

- For each binary many-to-many relationship set R



- Create a new relation schema S to represent R
- Include as foreign key attributes in S the primary keys of the relation schemas for the participating entity sets in R
- Their combination will form the primary key of S
- Also include attributes of the many-to-many relationship set as attributes of S






Map the many-to-many relationship sets

WORKS\_ON by creating the relation schema WORKS\_ON

Include the primary keys of PROJECT and EMPLOYEE as foreign keys

# Compare the following three choices to include WORKS\_ON



Add to EMPLOYEE

name	<u>id</u>	bdate	addr	sex	salary	dno	<u>pnumber</u>	hours
Yeung	7	080370	...	F	20K	3	Null	Null
Chan	3	031060	...	M	30K	4	C77	89
Chan	3	031060	...	M	30K	4	A01	10
Wong	4	010280	...	F	10K	7	A01	101
Cheung	8	220985	...	M	24K	1	A01	22
Cheung	8	220985	...	M	24K	1	C77	57

Add to PROJECT


pname	<u>pnumber</u>	plocation	dnum	<u>eid</u>	hours
SmartCard	C77	C-Lab	4	2	89
SmartCard	C77	C-Lab	4	8	10
Robotics	A01	C-Lab	7	2	101
Robotics	A01	C-Lab	7	4	22
Robotics	A01	C-Lab	7	8	57

New relation WORKS\_ON

<u>eid</u>	<u>pnumber</u>	hours
2	C77	89
2	A11	10
4	A11	101
8	A11	22
8	C77	57

# Step 6 (Non-binary Relationship)

- For each non-binary relationship set

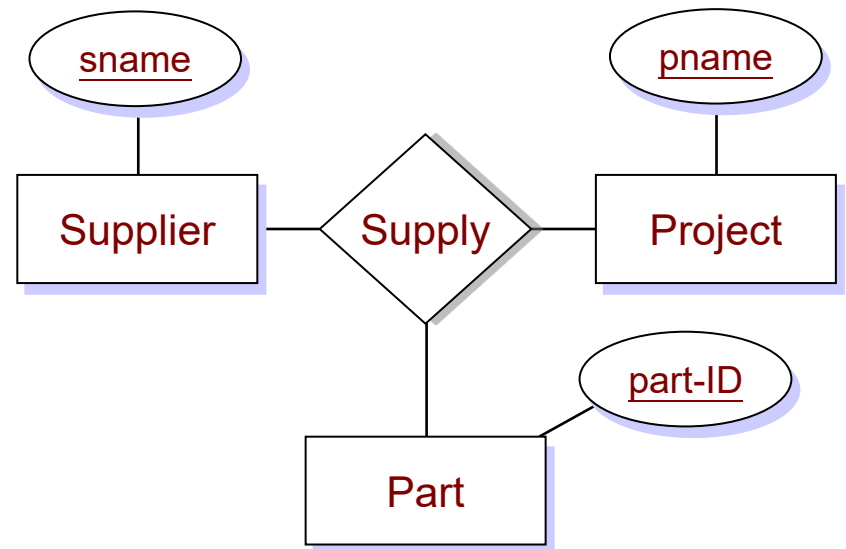



- Create a new relation schema S to represent R
- Include as foreign key attributes in S the primary keys of the participating entity sets
- Also include any attributes of the non-binary relationship set as attributes of S



# Ternary Relationship Set



- Create table Supply,
- Attributes are sname, pname, part-ID,
- These also form the key



Supply (sname, pname, part-ID)

# Ternary Relationship Set

- Create a new table X,
- Attributes are Case-ID, cname, lname,
- Case-ID is the key



$X(\underline{\text{Case-ID}}, \text{cname}, \text{lname})$

# Resulting Relation Schemas:

DEPARTMENT	dname	<u>dnumber</u>	mgr_id	mgr_start_date
------------	-------	----------------	--------	----------------

EMPLOYEE	name	<u>eid</u>	bdate	addr	sex	salary	dno	super_id
----------	------	------------	-------	------	-----	--------	-----	----------

PROJECT	pname	<u>pnumber</u>	plocation	dnum
---------	-------	----------------	-----------	------

DEPENDENT	<u>eid</u>	<u>dependent-name</u>	sex	bdate	relationship
-----------	------------	-----------------------	-----	-------	--------------

WORKS_ON	<u>eid</u>	<u>pnumber</u>	hours
----------	------------	----------------	-------

# Summary of Mapping Constructs

## Correspondence between ER and Relational Models

### ER Model

Entity type

1:1 or 1:N relationship type

M:N relationship type

*n*-ary relationship type

Simple attribute

Composite attribute

Multivalued attribute

Value set

Key attribute

### Relational Model

“Entity” relation

Foreign key (or “relationship” relation)

“Relationship” relation and two foreign keys

“Relationship” relation and *n* foreign keys

Attribute

Set of simple component attributes

Relation and foreign key

Domain

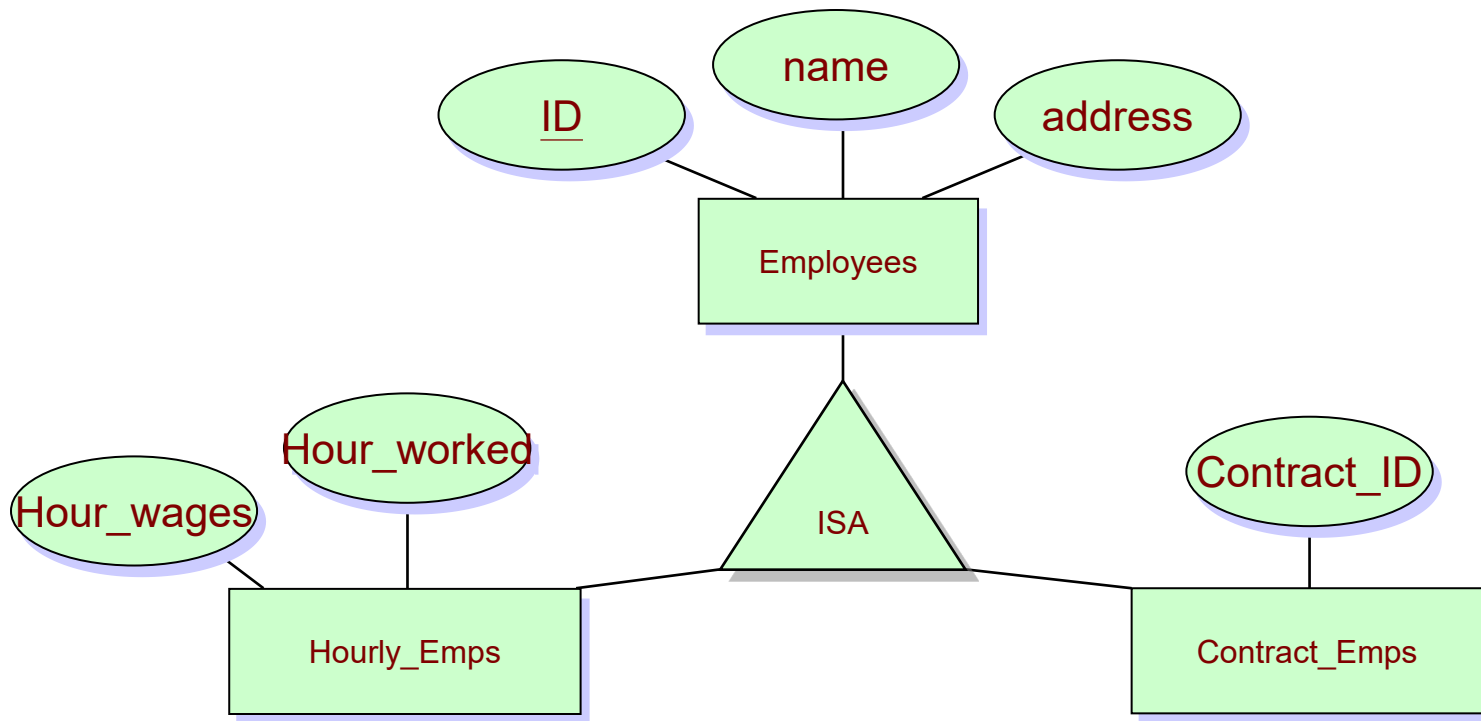
Primary (or secondary) key

# Outline

- Relational Model
  - Relational Model Concepts
  - Relational Database Schemas
- **ER-to-Relational Mapping**
  - Translating traditional ER diagrams
  - **Translating Class Hierarchy**

# Translating Class Hierarchy

- Consider the class hierarchy example



# Translating Class Hierarchy

- Two possible ways:
  - 1. Map each of the entity sets Employees, Hourly-Emps, and Contract-Emps to a distinct relation

EMPLOYEE	name	<u>id</u>	address
----------	------	-----------	---------

HOURLY-EMP	<u>id</u>	hours_worked	hours_wages
------------	-----------	--------------	-------------

CONTRACT-EMP	<u>id</u>	contract-id
--------------	-----------	-------------

# Translating Class Hierarchy

- Two possible ways:
  - 2. Create only two relations

HOURLY-EMP

<u>id</u>	name	address	hours_worked	hours_wages
-----------	------	---------	--------------	-------------

CONTRACT-EMP

<u>id</u>	name	address	contract-id
-----------	------	---------	-------------

- This requires the covering constraint to hold. (i.e. Hourly\_Emp and Contract\_Emp COVER Employee)

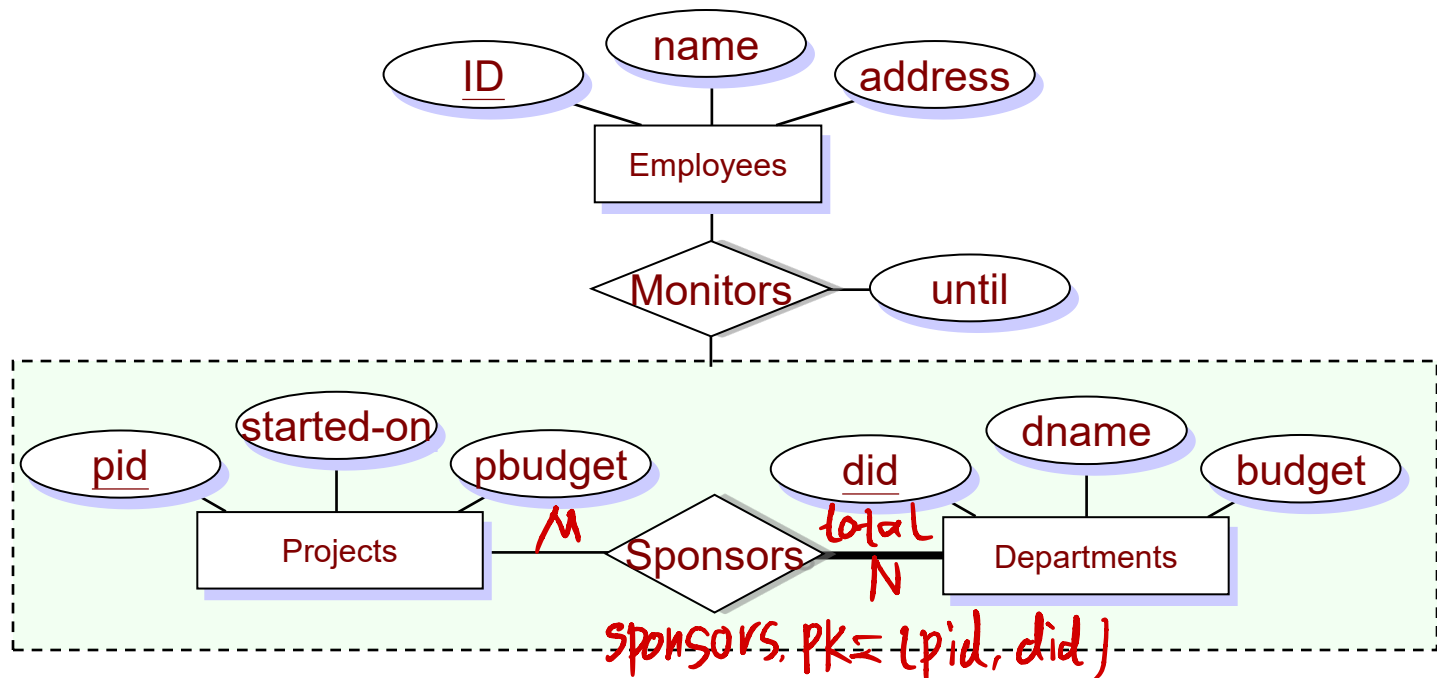


# Translating Class Hierarchy

- The first method is more general
- Disadvantage:
  - an extra relation is needed
  - more operation may be necessary when we need to get the employee information (e.g. looking up two relations)
- The second method is not always possible
- Advantage:
  - information of each employee is more easily accessible (usually only one relation look up)
  - however, if an employee can be both hourly and contract based, then information of the employee will be stored twice

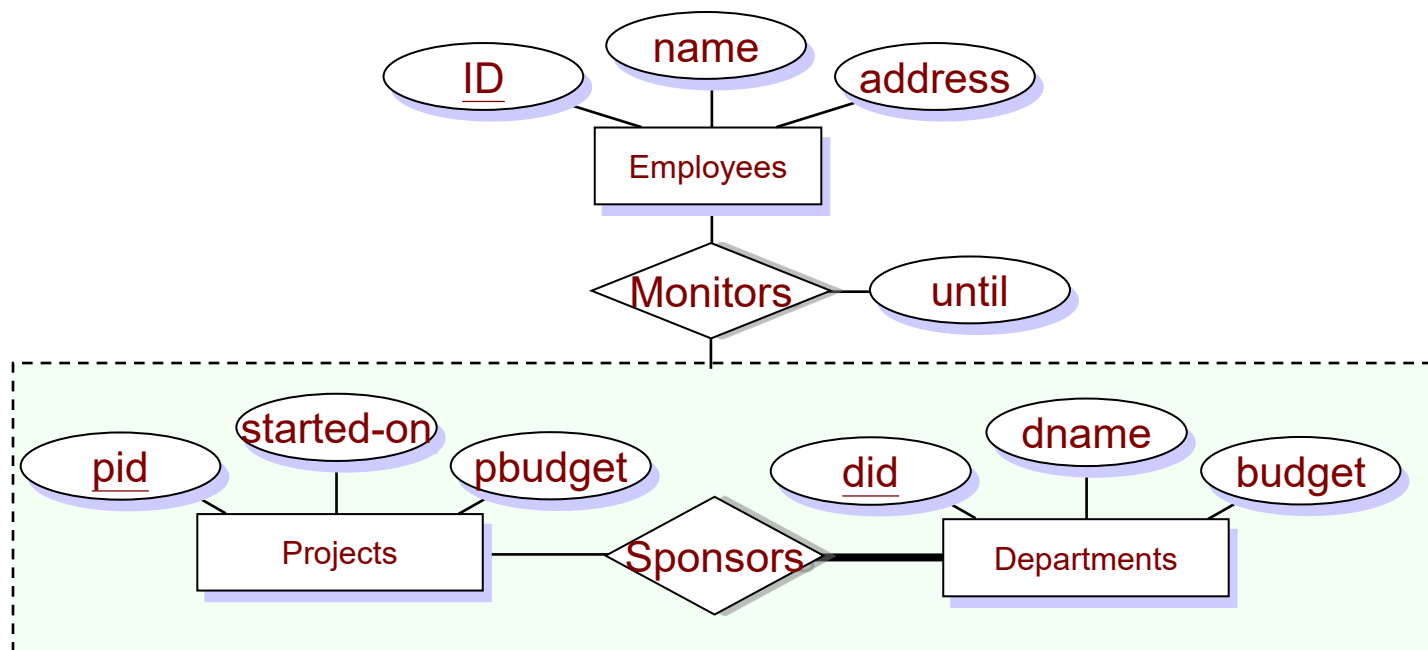
# Translating Aggregation

- There is no real distinction between entities and relationships in the relational model
- Therefore the mapping of aggregation in the E-R model to the relational model is quite simple
- For example,

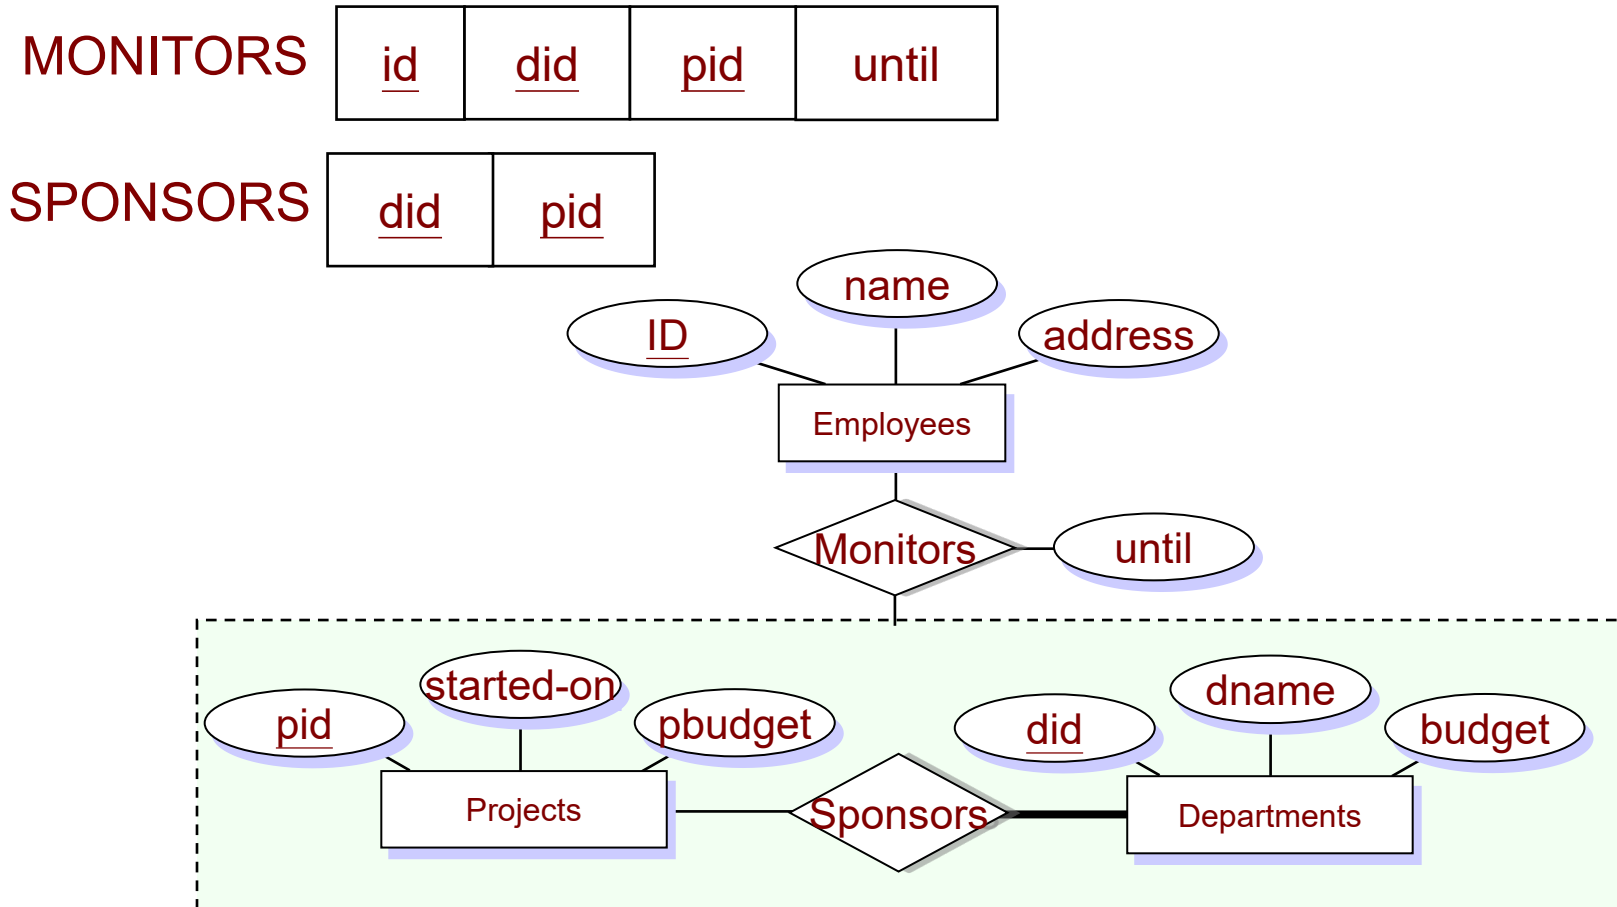


# Translating Aggregation

- For example,
  - The key attributes of SPONSORS are: (did, pid)
  - (a sponsorship is determined by department id and project id)



- The key attributes of SPONSORS are: (did, pid)
- For the Monitors relationship we create a relation:



- Sponsors is not contained in Monitors above, if a sponsorship has no monitor, then it will not appear in Monitors

# Summary

- What is a relation, relation schema
- How SQL creates tables (relations)
- Integrity Constraints - Primary Keys, Foreign Keys  
Referential Integrity
- Translating ER to relational model
  - Translating class hierarchies, aggregations
  - Reduce data redundancy
  - Reduce the number of NULL values
  - Reduce the cost of lookup for the information

# In-Class Exercise

Consider the following relations for a database that keeps track of student enrollment in courses and the books adopted for each course:

- STUDENT(SSN, Name, Major, Bdate)
- COURSE(Course#, Cname, Dept)
- ENROLL(SSN, Course#, Quarter, Grade)
- BOOK\_ADOPTION(Course#, Quarter, Book\_ISBN)
- TEXT(Book\_ISBN, Book\_Title, Publisher, Author)

**Draw a relational schema diagram specifying the foreign keys for this schema.**

